

## Lecture 6: theoretical foundations of concurrency

timeline, events, precedence, 2-thread mutual exclusion, deadlock freedom, starvation freedom, N-thread mutual exclusion, sequential objects and specifications, concurrent objects, linearizability

Alexander Filatov  
filatovaur@gmail.com

<https://github.com/Svazars/parallel-programming/blob/main/slides/pdf/16.pdf>

## In previous episodes

Concurrency is complicated due to:

- Independence of agents
  - combinatorial explosion of possible executions
- Uncontrollable execution speed of agents
  - spurious failures
- Uncertainty with definition of consistency
  - concurrent invariants
  - visibility
  - data race
  - atomicity

## In previous episodes

Concurrency is complicated due to:

- Independence of agents
  - combinatorial explosion of possible executions
- Uncontrollable execution speed of agents
  - spurious failures
- Uncertainty with definition of consistency
  - concurrent invariants
  - visibility
  - data race
  - atomicity

Is it purely engineering issue or we have encountered some fundamental problem?

## Be patient

Reminder: you are studying Turing Machines or Finite-State Machines in order to

- get a compact and fully formalized description of some "computing device"
- understand theoretical limitations of an object (Halting Problem, Pumping Lemma)

## Be patient

Reminder: you are studying Turing Machines or Finite-State Machines in order to

- get a compact and fully formalized description of some "computing device"
- understand theoretical limitations of an object (Halting Problem, Pumping Lemma)

It usually takes several mathematically-intensive sessions to go from the definition to really interesting theorems.

## Be patient

Reminder: you are studying Turing Machines or Finite-State Machines in order to

- get a compact and fully formalized description of some "computing device"
- understand theoretical limitations of an object (Halting Problem, Pumping Lemma)

It usually takes several mathematically-intensive sessions to go from the definition to really interesting theorems.

Concurrency is more "inherently-complicated" so it takes even more time to rigorously prove non-trivial properties.

## Be patient

Reminder: you are studying Turing Machines or Finite-State Machines in order to

- get a compact and fully formalized description of some "computing device"
- understand theoretical limitations of an object (Halting Problem, Pumping Lemma)

It usually takes several mathematically-intensive sessions to go from the definition to really interesting theorems.

Concurrency is more "inherently-complicated" so it takes even more time to rigorously prove non-trivial properties.

We will skip or simplify some steps.

## Be patient

Reminder: you are studying Turing Machines or Finite-State Machines in order to

- get a compact and fully formalized description of some "computing device"
- understand theoretical limitations of an object (Halting Problem, Pumping Lemma)

It usually takes several mathematically-intensive sessions to go from the definition to really interesting theorems.

Concurrency is more "inherently-complicated" so it takes even more time to rigorously prove non-trivial properties.

We will skip or simplify some steps.

There still will be a lot of technical (yet mathematically heavy) parts.

## Be patient

Reminder: you are studying Turing Machines or Finite-State Machines in order to

- get a compact and fully formalized description of some "computing device"
- understand theoretical limitations of an object (Halting Problem, Pumping Lemma)

It usually takes several mathematically-intensive sessions to go from the definition to really interesting theorems.

Concurrency is more "inherently-complicated" so it takes even more time to rigorously prove non-trivial properties.

We will skip or simplify some steps.

There still will be a lot of technical (yet mathematically heavy) parts.

At least, you will understand why proving concurrent algorithms could be really hard.

## Supplementary materials

- Chapters 1-5 in "The Art of Multiprocessor Programming"
- Companion materials <https://booksite.elsevier.com/9780123973375>

Our course briefly introduces some important theoretical concepts, for better understanding consider looking through

- chapter\_02.ppt
- chapter\_03.ppt

## Theoretical mode enabled

**Important:** in Lecture 6 (today) and Lecture 7 (next time) we will use **pseudocode**.

## Theoretical mode enabled

**Important:** in Lecture 6 (today) and Lecture 7 (next time) we will use **pseudocode**. It looks like Java and feels like Java but algorithms from the slides **will not work** as expected if run on the JVM.

## Theoretical mode enabled

**Important:** in Lecture 6 (today) and Lecture 7 (next time) we will use **pseudocode**.  
It looks like Java and feels like Java but algorithms from the slides **will not work** as expected if run on the JVM.  
In Lecture 9 and Lecture 10 we will understand why.

## Theoretical mode enabled

**Important:** in Lecture 6 (today) and Lecture 7 (next time) we will use **pseudocode**.  
It looks like Java and feels like Java but algorithms from the slides **will not work** as expected if run on the JVM.

In Lecture 9 and Lecture 10 we will understand why.

We are going to do some real Computer Science, it should work from mathematical point of view ;)

## Theoretical mode enabled

**Important:** in Lecture 6 (today) and Lecture 7 (next time) we will use **pseudocode**. It looks like Java and feels like Java but algorithms from the slides **will not work** as expected if run on the JVM.

In Lecture 9 and Lecture 10 we will understand why.

We are going to do some real Computer Science, it should work from mathematical point of view ;)

Theoretical intuition and insights will help us to understand how to implement efficient and correct algorithms using existing hardware.

## Theoretical mode enabled

**Important:** in Lecture 6 (today) and Lecture 7 (next time) we will use **pseudocode**. It looks like Java and feels like Java but algorithms from the slides **will not work** as expected if run on the JVM.

In Lecture 9 and Lecture 10 we will understand why.

We are going to do some real Computer Science, it should work from mathematical point of view ;)

Theoretical intuition and insights will help us to understand how to implement efficient and correct algorithms using existing hardware.

**Hint:** focus on ideas rather mathematical formalisms.

# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - Intuition behind consistency definitions
  - Optional: formal definitions and composability
- 4 Summary

# Time and events

An event  $a_0$  of thread A is

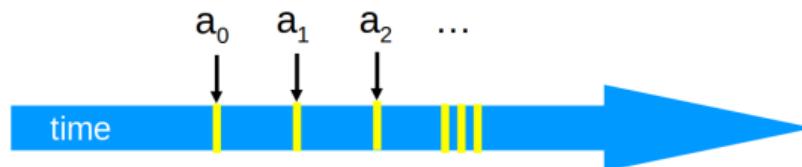
- Instantaneous
- No simultaneous events (break ties)



# Threads

A thread A is (formally) a sequence  $a_0, a_1, \dots$  of events

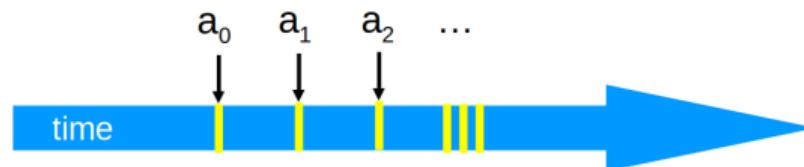
- “Trace” model
- Notation:  $a_0 \rightarrow a_1$  indicates order



# Threads

A thread A is (formally) a sequence  $a_0, a_1, \dots$  of events

- “Trace” model
- Notation:  $a_0 \rightarrow a_1$  indicates order

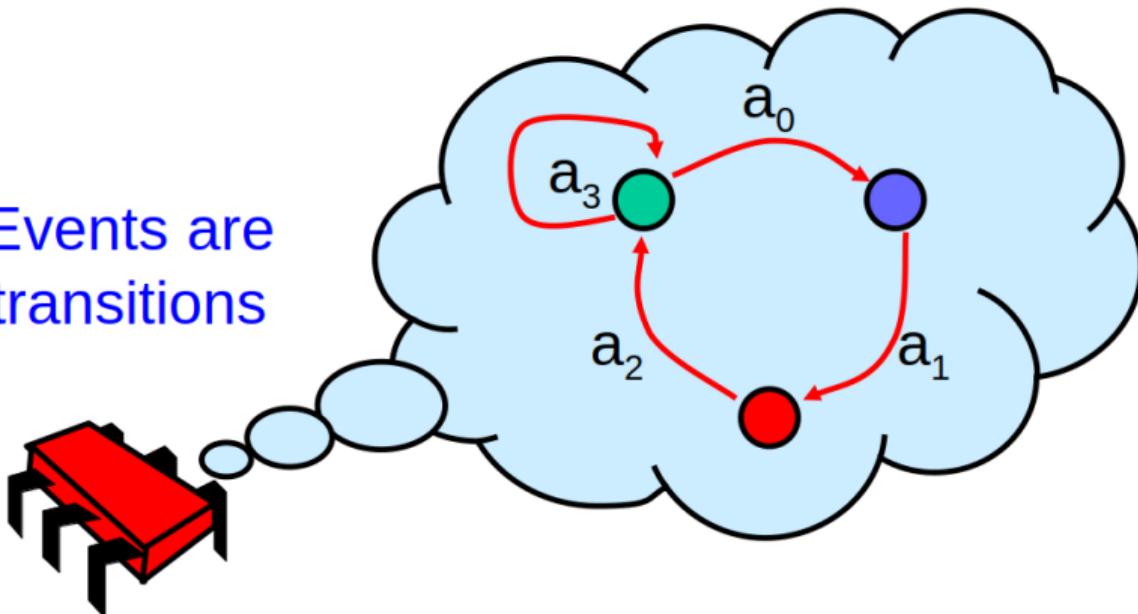


Example Thread events:

- Assign to shared or local variable
- Invoke method
- Return from method
- ...

# Threads are state machines

Events are transitions



# Concurrent threads

Thread A



Thread B



Events of two or more threads:

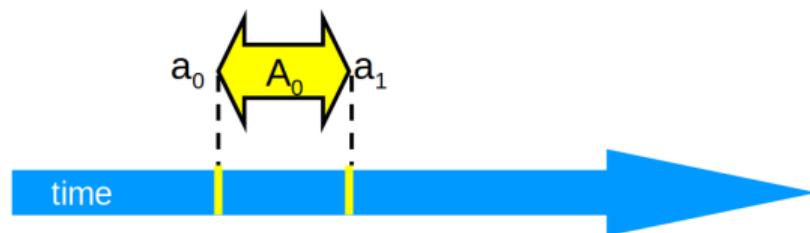
- Interleaved
- Not necessarily independent



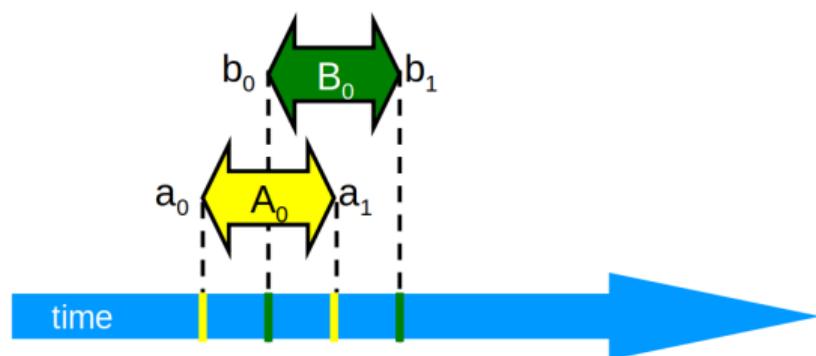
# Methods are intervals

An **interval**  $A_0 = (a_0, a_1)$  is

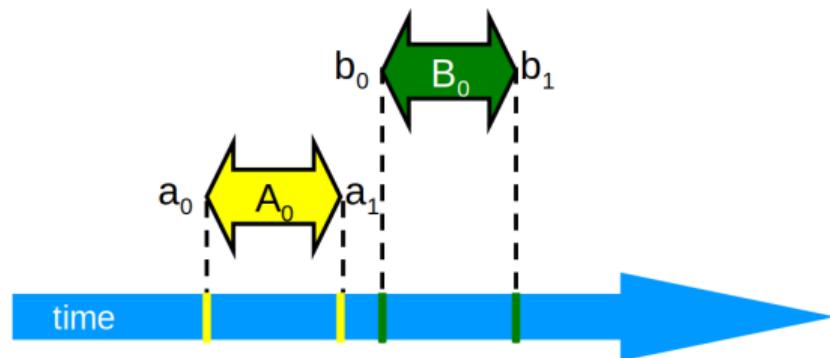
- Time between events  $a_0$  and  $a_1$



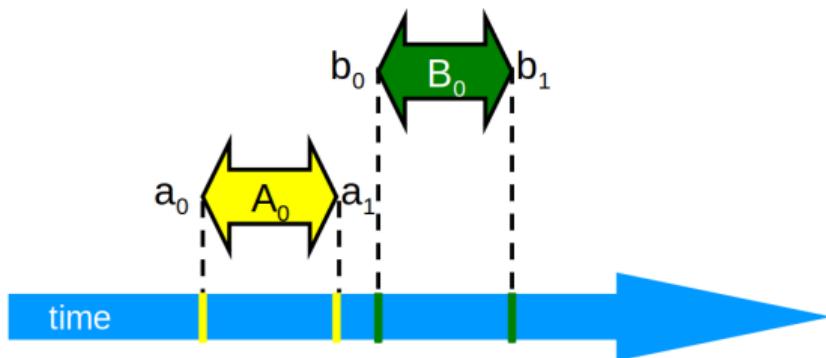
# Overlapping intervals



# Disjoint intervals



# Disjoint intervals



**Precedence:**

- Interval  $A_0$  precedes interval  $B_0$
- Notation:  $A_0 \rightarrow B_0$

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

:

- Never true that  $A \rightarrow A$ .
- If  $A \rightarrow B$  then not true  $B \rightarrow A$ .
- If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ .

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

:

- Never true that  $A \rightarrow A$ .
- If  $A \rightarrow B$  then not true  $B \rightarrow A$ .
- If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ .
- Remember:  $A \rightarrow B$  and  $B \rightarrow A$  might both be false!

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

Precedence is a partial order:

- Never true that  $A \rightarrow A$ .
- If  $A \rightarrow B$  then not true  $B \rightarrow A$ .
- If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ .
- Remember:  $A \rightarrow B$  and  $B \rightarrow A$  might both be false!

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

Precedence is a partial order:

- Never true that  $A \rightarrow A$ . (Irreflexive)
- If  $A \rightarrow B$  then not true  $B \rightarrow A$ .
- If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ .
- Remember:  $A \rightarrow B$  and  $B \rightarrow A$  might both be false!

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

Precedence is a partial order:

- Never true that  $A \rightarrow A$ . (Irreflexive)
- If  $A \rightarrow B$  then not true  $B \rightarrow A$ . (Antisymmetric)
- If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ .
- Remember:  $A \rightarrow B$  and  $B \rightarrow A$  might both be false!

# Precedence

$$A_0 \rightarrow B_0$$

- $A_0$  precedes interval  $B_0$
- $A_0$  happens before interval  $B_0$
- Fancy way to say "Middle Ages → Renaissance"

what about *this week* vs *this month*?

Precedence is a partial order:

- Never true that  $A \rightarrow A$ . (Irreflexive)
- If  $A \rightarrow B$  then not true  $B \rightarrow A$ . (Antisymmetric)
- If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ . (Transitive)
- Remember:  $A \rightarrow B$  and  $B \rightarrow A$  might both be false!

# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - Intuition behind consistency definitions
  - Optional: formal definitions and composability
- 4 Summary

# Mutual exclusion

Thread is *well formed* if:

- each critical section is associated with a unique Lock object
- Lock.lock when trying to enter critical section
- Lock.unlock when leaving critical section

Let  $CS_A^j$  be the interval: thread A executes critical section for j-th time.

## Mutual exclusion for 2 agents

$CS_A^j$ : thread A executes critical section for j-th time.

## Mutual exclusion for 2 agents

$CS_A^j$ : thread A executes critical section for j-th time.

**Mutual exclusion:**

- Critical sections of different threads do not overlap.

## Mutual exclusion for 2 agents

$CS_A^j$ : thread A executes critical section for j-th time.

**Mutual exclusion:**

- Critical sections of different threads do not overlap.  $\forall A, B, k, j \ CS_A^k \rightarrow CS_B^j \text{ or } CS_B^j \rightarrow CS_A^k$

## Mutual exclusion for 2 agents

$CS_A^j$ : thread A executes critical section for j-th time.

### Mutual exclusion:

- Critical sections of different threads do not overlap.  $\forall A, B, k, j \ CS_A^k \rightarrow CS_B^j \text{ or } CS_B^j \rightarrow CS_A^k$

### Freedom from deadlock:

- If some thread attempts to acquire the lock, then some thread will succeed in acquiring the lock. If thread A calls `lock()` but never acquires the lock, then other threads must be completing an infinite number of critical sections.

## Mutual exclusion for 2 agents

$CS_A^j$ : thread A executes critical section for j-th time.

### Mutual exclusion:

- Critical sections of different threads do not overlap.  $\forall A, B, k, j \ CS_A^k \rightarrow CS_B^j \text{ or } CS_B^j \rightarrow CS_A^k$

### Freedom from deadlock:

- If some thread attempts to acquire the lock, then some thread will succeed in acquiring the lock. If thread A calls `lock()` but never acquires the lock, then other threads must be completing an infinite number of critical sections.

### Freedom from starvation (lockout freedom):

- Every thread that attempts to acquire the lock eventually succeeds. Every call to `lock()` eventually returns.

## Mutual exclusion for 2 agents

$CS_A^j$ : thread A executes critical section for j-th time.

### Mutual exclusion:

- Critical sections of different threads do not overlap.  $\forall A, B, k, j \ CS_A^k \rightarrow CS_B^j \text{ or } CS_B^j \rightarrow CS_A^k$

### Freedom from deadlock:

- If some thread attempts to acquire the lock, then some thread will succeed in acquiring the lock. If thread A calls `lock()` but never acquires the lock, then other threads must be completing an infinite number of critical sections.

### Freedom from starvation (lockout freedom):

- Every thread that attempts to acquire the lock eventually succeeds. Every call to `lock()` eventually returns.

### Theorem

*Starvation freedom implies deadlock freedom.*

## Question time

Question: Mutual exclusion, Freedom from deadlock, Freedom from starvation. Which are safety properties, which are liveness properties?



## Question time

Question: Restaurant "Concurrency pitfalls" claims that every philosopher will eat ordered meal before Sun explodes. Does it provide **Freedom from starvation?**



## LockOne

```
class LockOne implements Lock {  
    boolean[] flag = new boolean[2];  
    void lock() {  
        int i = ThreadID.get(); // 0 or 1  
        int j = 1 - i;  
        flag[i] = true;          // I want to enter  
        while (flag[j]) {}      // await other thread  
    }  
    void unlock() {  
        int i = ThreadID.get();  
        flag[i] = false;         // I do not want to enter any more  
    }  
}
```

## LockOne: deadlock

```
void lock() {  
    int i = ThreadID.get();  
    int j = 1 - i;  
    flag[i] = true;  
    while (flag[j]) {}  
}  
void unlock() {  
    flag[ThreadID.get()] = false;  
}
```

$\text{write}_A(\text{flag}[A] = \text{true})$  occur before  $\text{read}_B(\text{flag}[A])$   
 $\text{write}_B(\text{flag}[B] = \text{true})$  occur before  $\text{read}_A(\text{flag}[B])$

## LockOne: mutual exclusion

```
void lock() {  
    flag[i] = true;  
    while (flag[j]) {}  
}  
void unlock() {  
    flag[i] = false;  
}
```

From the code:

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] = \text{false}) \rightarrow CS_A$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] = \text{false}) \rightarrow CS_B$

## LockOne: mutual exclusion

Plan:

- Assume  $CS_A^j$  overlaps  $CS_B^k$ .
- Consider each thread's last ( $j$ -th and  $k$ -th) read and write in the `lock()` method before entering
- Derive a contradiction

## LockOne: mutual exclusion

Plan:

- Assume  $CS_A^j$  overlaps  $CS_B^k$ .
- Consider each thread's last ( $j$ -th and  $k$ -th) read and write in the `lock()` method before entering
- Derive a contradiction

Informal hint: Thread A observed `flag[B] = false` and entered  $CS_A^j$  and then Thread B entered  $CS_B^k$  which required to write `flag[B] = true`

## LockOne: mutual exclusion

Plan:

- Assume  $CS_A^j$  overlaps  $CS_B^k$ .
- Consider each thread's last ( $j$ -th and  $k$ -th) read and write in the `lock()` method before entering
- Derive a contradiction

Informal hint: Thread A observed `flag[B] = false` and entered  $CS_A^j$  and then Thread B entered  $CS_B^k$  which required to write `flag[B] = true`

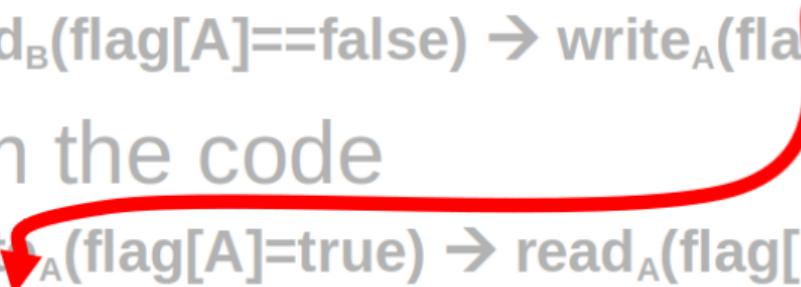
From assumption:

- $read_A(flag[B] = \text{false}) \rightarrow write_B(flag[B] = \text{true})$
- $read_B(flag[A] = \text{false}) \rightarrow write_A(flag[A] = \text{true})$

## LockOne: mutual exclusion

- Assumptions:
  - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
  - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
  - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
  - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## LockOne: mutual exclusion

- Assumptions:
    - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
    - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
  - From the code
    - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
    - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 

## LockOne: mutual exclusion

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## LockOne: mutual exclusion

- Assumptions:
  - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
  - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
  - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
  - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

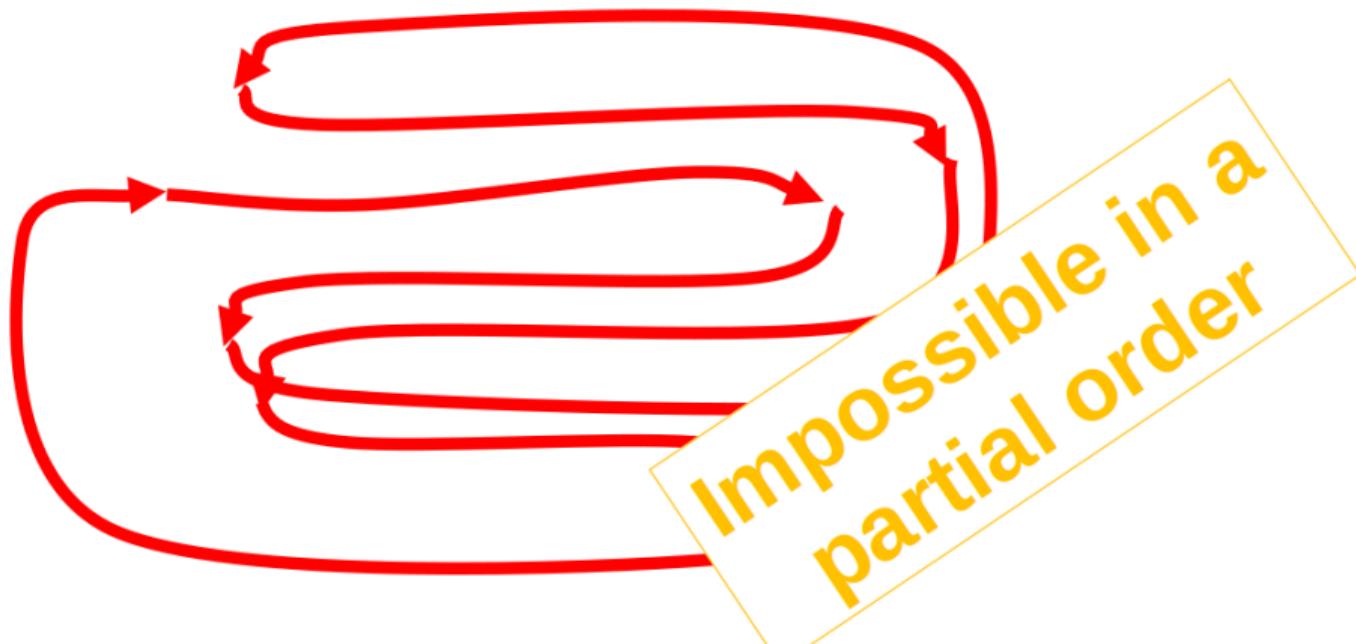
## LockOne: mutual exclusion

- Assumptions.
  - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
  - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
  - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
  - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## LockOne: mutual exclusion

- Assumptions.
  - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
  - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
  - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
  - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## LockOne: mutual exclusion



# LockOne

## Properties

- Mutual exclusion
- Deadlock-freedom

# LockOne

## Properties

- Mutual exclusion
- Deadlock-freedom

```
flag[i] = true;  
while (flag[j]) {}
```

```
flag[j] = true;  
while (flag[i]) {}
```

# LockOne

## Properties

- Mutual exclusion
- Deadlock-freedom

```
flag[i] = true;           flag[j] = true;  
while (flag[j]) {}        while (flag[i]) {}
```

- Concurrent executions can deadlock
- Sequential executions OK

## LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get(); // 0 or 1  
        victim = i; // let the other go first  
        while (victim == i) {}; // wait  
    }  
    public void unlock() {}  
}
```

## LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get(); // 0 or 1  
        victim = i; // let the other go first  
        while (victim == i) {}; // wait  
    }  
    public void unlock() {}  
}
```

Satisfies Mutual Exclusion:

- if thread  $i$  in  $CS$
- then  $victim == j$
- cannot be both 0 and 1

## LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get(); // 0 or 1  
        victim = i;           // let the other go first  
        while (victim == i) {}; // wait  
    }  
    public void unlock() {}  
}
```

## LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get(); // 0 or 1  
        victim = i;           // let the other go first  
        while (victim == i) {}; // wait  
    }  
    public void unlock() {}  
}
```

Not deadlock free:

- Sequential execution deadlocks
- Concurrent execution does not

# LockOne and LockTwo

LockOne, LockTwo:

- Mutual exclusion
- Deadlock-freedom

# LockOne and LockTwo

LockOne, LockTwo:

- Mutual exclusion
- Deadlock-freedom

LockOne:

- Concurrent executions can deadlock
- Sequential executions OK

# LockOne and LockTwo

LockOne, LockTwo:

- Mutual exclusion
- Deadlock-freedom

LockOne:

- Concurrent executions can deadlock
- Sequential executions OK

LockTwo:

- Concurrent execution OK
- Sequential execution can deadlock

# LockOne and LockTwo

LockOne, LockTwo:

- Mutual exclusion
- Deadlock-freedom

LockOne:

- Concurrent executions can deadlock
- Sequential executions OK

LockTwo:

- Concurrent execution OK
- Sequential execution can deadlock

Idea: combine them!

## Peterson's Algorithm

```
class Peterson implements Lock {  
    private boolean[] flag = new boolean[2];  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get(); // 0 or 1  
        int j = 1 - i;  
        flag[i] = true;           // I am interested  
        victim = i;              // you go first  
        while (flag[j] && victim == i) {}; // wait while other interested and  
                                         // I am the victim  
    }  
    public void unlock() {  
        flag[ThreadID.get()] = false; // I am not interested  
    }  
}
```

# Peterson's Algorithm

## Mutual exclusion

```
public void lock() {  
    flag[i] = true;          // I am interested  
    victim = i;              // you go first  
    while (flag[j] && victim == i) {}; // wait while other interested and  
                                         // I am the victim  
}
```

From the code:

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{write}_B(\text{victim} = B)$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$

Assumption (both threads entered CS, latest to write victim is thread A):

- $\text{write}_B(\text{victim} = B) \rightarrow \text{write}_A(\text{victim} = A)$

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{write}_B(\text{victim} = B)$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$
- $\text{write}_B(\text{victim} = B) \rightarrow \text{write}_A(\text{victim} = A)$

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{write}_B(\text{victim} = B)$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$
- $\text{write}_B(\text{victim} = B) \rightarrow \text{write}_A(\text{victim} = A)$

Let's swap second and third lines.

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{write}_B(\text{victim} = B)$
- $\text{write}_B(\text{victim} = B) \rightarrow \text{write}_A(\text{victim} = A)$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{write}_B(\text{victim} = B)$
- $\text{write}_B(\text{victim} = B) \rightarrow \text{write}_A(\text{victim} = A)$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$

Let's remove repeats.

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow$
- $\text{write}_B(\text{victim} = B) \rightarrow$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow$
- $\text{write}_B(\text{victim} = B) \rightarrow$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$

Thread A read `flag[B]` == true and `victim` == A so it could not have entered the CS.

# Peterson's Algorithm

## Mutual exclusion

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow$
- $\text{write}_B(\text{victim} = B) \rightarrow$
- $\text{write}_A(\text{victim} = A) \rightarrow \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim})$

Thread A read `flag[B] == true` and `victim == A` so it could not have entered the CS.  
QED.

# Peterson's Algorithm

Starvation freedom (implies deadlock freedom)

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

# Peterson's Algorithm

Starvation freedom (implies deadlock freedom)

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

- Thread  $i$  blocked only if  $j$  repeatedly re-enters so  $\text{flag}[j] == \text{true}$  and  $\text{victim} == i$ .

# Peterson's Algorithm

Starvation freedom (implies deadlock freedom)

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

- Thread i blocked only if j repeatedly re-enters so `flag[j] == true` and `victim == i`.
- When thread j re-enters it sets `victim = j`.

# Peterson's Algorithm

Starvation freedom (implies deadlock freedom)

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

- Thread i blocked only if j repeatedly re-enters so `flag[j] == true` and `victim == i`.
- When thread j re-enters it sets `victim = j`.
- So thread i gets in.

## Summary: 2-thread mutual exclusion

Formalize concurrent method execution:

- Single timeline
- Totally ordered atomic events
- Methods represented as partially ordered intervals
- *Precedence* is irreflexive, antisymmetric, transitive partial order

## Summary: 2-thread mutual exclusion

Formalize concurrent method execution:

- Single timeline
- Totally ordered atomic events
- Methods represented as partially ordered intervals
- *Precedence* is irreflexive, antisymmetric, transitive partial order

Formalize mutex properties:

- Mutual exclusion
- Deadlock freedom
- Starvation freedom (implies Deadlock freedom)

## Summary: 2-thread mutual exclusion

Formalize concurrent method execution:

- Single timeline
- Totally ordered atomic events
- Methods represented as partially ordered intervals
- *Precedence* is irreflexive, antisymmetric, transitive partial order

Formalize mutex properties:

- Mutual exclusion
- Deadlock freedom
- Starvation freedom (implies Deadlock freedom)

Prove theorems:

- Algorithm has property (*suppose not ...*)
- Algorithm lacks property (provide counterexample)

# Lecture plan

1 Time, events, intervals, precedence

2 Mutual exclusion

- 2 threads
- N threads
- Lower bounds on the Number of Locations
- Optional: fairness
- Mutual exclusion beyond read-write memory cells

3 Concurrent objects

- Sequential and concurrent objects
- Intuition behind consistency definitions
- Optional: formal definitions and composability

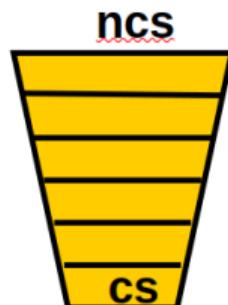
4 Summary

# Filter Lock

## Idea

There are  $n-1$  “waiting rooms” called levels.

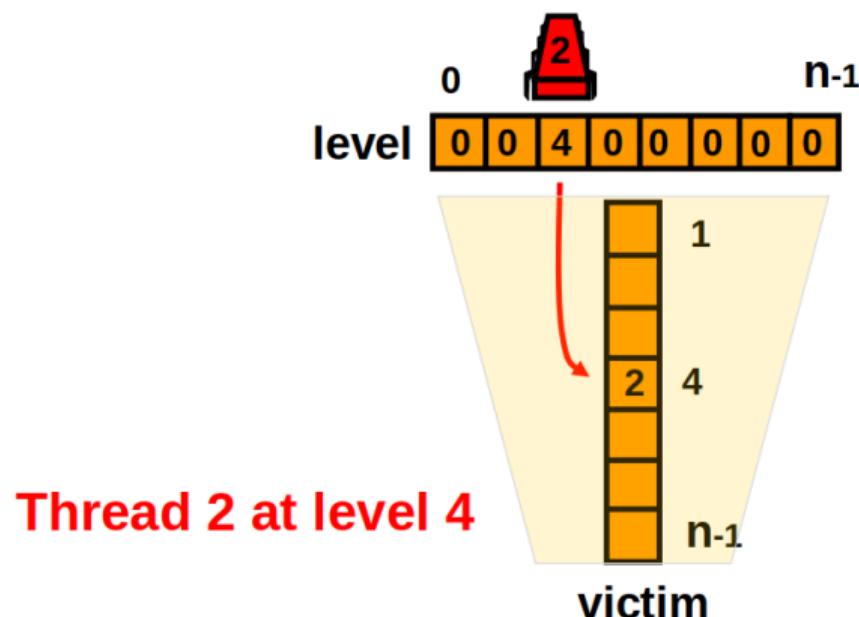
- At each level
  - At least one enters level
  - At least one blocked if many try
- Only one thread makes it through



# Filter Lock

## Level and victim

```
class Filter implements Lock {
    // level[i] for thread i
    int[] level;
    // victim[L] for level L
    int[] victim;
    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 1; i < n; i++) {
            level[i] = 0;
        }
    }
}
```



# Filter Lock

```
class Filter implements Lock {  
    ...  
  
    public void lock(){  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i level[k] >= L) &&  
                   victim[L] == i ) {};  
        }  
        public void unlock() {  
            level[i] = 0;  
        }  
    }  
}
```

# Filter Lock

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                  victim[L] == i) {};  
        }  
        public void release(int i) {  
            level[i] = 0;  
        }  
    }  
}
```

One level at a time

# Filter Lock

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                  victim[L] == i)  
        }  
        public void release(int i)  
            level[i] = 0;  
    }  
}
```

Announce  
intention to enter  
level L

# Filter Lock

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                  victim[L] == i) {};  
        }  
        public void release(int i)  
            level[i] = 0;  
    }  
}
```

Give priority to  
anyone but me

# Filter Lock

Wait as long as someone else is at same or higher level, and I'm designated victim

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists$  k != i) level[k] >= L) &&  
              victim[L] == i) {};  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

# Filter Lock

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                  victim[L] == i) {};  
        }  
    }  
}
```

Thread enters level L when it completes  
the loop

# Filter Lock

## Proof idea

- Start at level  $L = 0$
- At most  $n - L$  threads enter level  $L$
- Mutual exclusion at level  $L = n-1$

# Filter Lock

## Proof idea

- Start at level  $L = 0$
- At most  $n - L$  threads enter level  $L$
- Mutual exclusion at level  $L = n-1$

Induction Hypothesis:

- No more than  $n - (L-1)$  at level  $L-1$
- Induction step: by contradiction
- Assume all at level  $L-1$  enter level  $L$
- Thread A last to write  $\text{victim}[L]$
- Thread B is any other thread at level  $L$

# Filter Lock

## Proof idea

- Start at level  $L = 0$
- At most  $n - L$  threads enter level  $L$
- Mutual exclusion at level  $L = n-1$

Induction Hypothesis:

- No more than  $n - (L-1)$  at level  $L-1$
- Induction step: by contradiction
- Assume all at level  $L-1$  enter level  $L$
- Thread A last to write  $\text{victim}[L]$
- Thread B is any other thread at level  $L$

Homework: slides 77-93 from chapter\_02.ppt

Homework: Section 2.4 "The Filter Lock" (pages 28 - 31) from "The Art of Multiprocessor Programming"

# Filter Lock

No starvation

Filter Lock satisfies properties:

- Just like Peterson Alg at any level
- So no one starves

# Filter Lock

No starvation

Filter Lock satisfies properties:

- Just like Peterson Alg at any level
- So no one starves

But what about fairness?

# Filter Lock

No starvation

Filter Lock satisfies properties:

- Just like Peterson Alg at any level
- So no one starves

But what about fairness?

Threads can be overtaken by others

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

---

<sup>1</sup>Historical name

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

- **How many** registers are needed to solve N-thread mutual exclusion?

---

<sup>1</sup>Historical name

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

- **How many** registers are needed to solve N-thread mutual exclusion?
- **What** are those registers? How to describe them formally?

---

<sup>1</sup>Historical name

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

- **How many** registers are needed to solve N-thread mutual exclusion?
- **What** are those registers? How to describe them formally?
- Is there any **alternative** to registers and how to use it to implement mutual exclusion?

---

<sup>1</sup>Historical name

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

- **How many** registers are needed to solve N-thread mutual exclusion?  
Theorem about Lower bounds, next slide
- **What** are those registers? How to describe them formally?
- Is there any **alternative** to registers and how to use it to implement mutual exclusion?

---

<sup>1</sup>Historical name

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

- **How many** registers are needed to solve N-thread mutual exclusion?  
Theorem about Lower bounds, next slide
- **What** are those registers? How to describe them formally?  
Concurrent objects chapter, later today
- Is there any **alternative** to registers and how to use it to implement mutual exclusion?

---

<sup>1</sup>Historical name

## Summary: N-thread mutual exclusion

N-level design:

- every step "filters" (prevents from progress) at least one thread
- reuse ideas from simpler design (2-thread Peterson's algorithm)

We have many memory cells (**registers**<sup>1</sup>): boolean `flag[i]`, integer `victim[j]`

- **How many** registers are needed to solve N-thread mutual exclusion?  
Theorem about Lower bounds, next slide
- **What** are those registers? How to describe them formally?  
Concurrent objects chapter, later today
- Is there any **alternative** to registers and how to use it to implement mutual exclusion?  
Lecture 8

---

<sup>1</sup>Historical name

# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - Intuition behind consistency definitions
  - Optional: formal definitions and composability
- 4 Summary

# Mutual exclusion for $N$ threads requires read-write of at least $N$ locations

Optional homework:

*Theorem 2.8.1. Any Lock algorithm that, by reading and writing memory, solves deadlock-free mutual exclusion for three threads must use at least three distinct memory locations.*

Key insights:

- Covering state

Generalized theorem:  $N$ -thread deadlock-free mutual exclusion requires  $N$  distinct locations.

# Mutual exclusion: you know so far

Mutual exclusion could be solved:

- For  $N$  concurrent threads
- Using  $N$  boolean and  $N$  integer registers
- Provides starvation-freedom

# Mutual exclusion: you know so far

Mutual exclusion could be solved:

- For  $N$  concurrent threads
- Using  $N$  boolean and  $N$  integer registers
- Provides starvation-freedom

Mutual exclusion could not be solved:

- Using less than  $N$  registers

# Mutual exclusion: you know so far

Mutual exclusion could be solved:

- For  $N$  concurrent threads
- Using  $N$  boolean and  $N$  integer registers
- Provides starvation-freedom

Mutual exclusion could not be solved:

- Using less than  $N$  registers

Few questions remain:

## Mutual exclusion: you know so far

Mutual exclusion could be solved:

- For  $N$  concurrent threads
- Using  $N$  boolean and  $N$  integer registers
- Provides starvation-freedom

Mutual exclusion could not be solved:

- Using less than  $N$  registers

Few questions remain:

- How much time will thread wait until success? FilterLock does not guarantee fairness

## Mutual exclusion: you know so far

Mutual exclusion could be solved:

- For  $N$  concurrent threads
- Using  $N$  boolean and  $N$  integer registers
- Provides starvation-freedom

Mutual exclusion could not be solved:

- Using less than  $N$  registers

Few questions remain:

- How much time will thread wait until success? FilterLock does not guarantee fairness
- Could we "improve" registers with some additional operations to speed-up the algorithm?

## Mutual exclusion: you know so far

Mutual exclusion could be solved:

- For  $N$  concurrent threads
- Using  $N$  boolean and  $N$  integer registers
- Provides starvation-freedom

Mutual exclusion could not be solved:

- Using less than  $N$  registers

Few questions remain:

- How much time will thread wait until success? FilterLock does not guarantee fairness
- Could we "improve" registers with some additional operations to speed-up the algorithm?
- What if threads are created and deleted dynamically?

# Lecture plan

1 Time, events, intervals, precedence

2 Mutual exclusion

- 2 threads
- N threads
- Lower bounds on the Number of Locations
- **Optional: fairness**
- Mutual exclusion beyond read-write memory cells

3 Concurrent objects

- Sequential and concurrent objects
- Intuition behind consistency definitions
- Optional: formal definitions and composability

4 Summary

# Formalizing fairness

Optional homework: slides 95-99 from chapter\_02.ppt

Optional homework: Section 2.5 "Fairness" (page 31)

Key insights:

- *doorway* section
- *waiting* section
- first-come-first-served ( $D_A^j \rightarrow D_B^k$  then  $CS_A^j \rightarrow CS_B^k$ )

# Bakery algorithm

Optional homework: slides 99-118 from chapter\_02.ppt

Optional homework: Section 2.6 "Lamport's Bakery Algorithm" (pages 31 - 33)

Key insights:

- monotonically growing labels for operations
- proving fairness

# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - Intuition behind consistency definitions
  - Optional: formal definitions and composability
- 4 Summary

## Mutual exclusion using read/write locations

- Correct, starvation-free algorithm exists (filter lock)
- Correct, starvation-free and fair algorithm exists (bakery algorithm)
- Any correct algorithm will require at least  $N$  registers with read and write operations

## Mutual exclusion using read/write locations

- Correct, starvation-free algorithm exists (filter lock)
- Correct, starvation-free and fair algorithm exists (bakery algorithm)
- Any correct algorithm will require at least  $N$  registers with read and write operations

We could classify registers by allowed concurrency:

- Single Reader Single Writer (SRSW)
- Multiple Reader Single Writer (MRSW)
- Multiple Reader Multiple Writer (MRMW)

## Mutual exclusion using read/write locations

- Correct, starvation-free algorithm exists (filter lock)
- Correct, starvation-free and fair algorithm exists (bakery algorithm)
- Any correct algorithm will require at least  $N$  registers with read and write operations

We could classify registers by allowed concurrency:

- Single Reader Single Writer (SRSW)
- Multiple Reader Single Writer (MRSW)
- Multiple Reader Multiple Writer (MRMW)

We could classify registers by capacity:

- Boolean register (*true/false*), Integer register ( $0, 1, \dots, 2^N$ )

## Mutual exclusion using read/write locations

- Correct, starvation-free algorithm exists (filter lock)
- Correct, starvation-free and fair algorithm exists (bakery algorithm)
- Any correct algorithm will require at least  $N$  registers with read and write operations

We could classify registers by allowed concurrency:

- Single Reader Single Writer (SRSW)
- Multiple Reader Single Writer (MRSW)
- Multiple Reader Multiple Writer (MRMW)

We could classify registers by capacity:

- Boolean register (*true/false*), Integer register ( $0, 1, \dots, 2^N$ )

We could classify registers by supported operations:

- Read-only, Write-once, Monotonically increasing, Arbitrary update

## Mutual exclusion using read/write locations

- Correct, starvation-free algorithm exists (filter lock)
- Correct, starvation-free and fair algorithm exists (bakery algorithm)
- Any correct algorithm will require at least  $N$  registers with read and write operations

We could classify registers by allowed concurrency:

- Single Reader Single Writer (SRSW)
- Multiple Reader Single Writer (MRSW)
- Multiple Reader Multiple Writer (MRMW)

We could classify registers by capacity:

- Boolean register (*true/false*), Integer register ( $0, 1, \dots, 2^N$ )

We could classify registers by supported operations:

- Read-only, Write-once, Monotonically increasing, Arbitrary update
- Atomic read-modify-write operations

## Mutual exclusion using read/write locations

- Correct, starvation-free algorithm exists (filter lock)
- Correct, starvation-free and fair algorithm exists (bakery algorithm)
- Any correct algorithm will require at least  $N$  registers with read and write operations

We could classify registers by allowed concurrency:

- Single Reader Single Writer (SRSW)
- Multiple Reader Single Writer (MRSW)
- Multiple Reader Multiple Writer (MRMW)

We could classify registers by capacity:

- Boolean register (*true/false*), Integer register ( $0, 1, \dots, 2^N$ )

We could classify registers by supported operations:

- Read-only, Write-once, Monotonically increasing, Arbitrary update
- Atomic read-modify-write operations

Before diving into classifications (Lecture 7, Lecture 8), let's study few more mathematical instruments.

## Homework: before final exam

When you finish all lectures of second block, you will know many advanced concepts:

- Concurrent object consistency definitions
- Register constructions
- Atomic snapshots
- Consensus number

Revise all the correctness proofs from this lecture and be ready to answer some tricky questions:

- In order to prove correctness of Peterson's algorithm, should we use MRSW or MRMW register?
- In order to prove correctness of mutual exclusion primitives, should we require all registers to be linearizable or weaker guarantees will fit?
- Could you adapt FilterLock to allow dynamic thread creation and deletion?

# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - Intuition behind consistency definitions
  - Optional: formal definitions and composability
- 4 Summary

# Motivation

We need to formalize

- Safety (correctness)
- Liveness (progress)

# Motivation

We need to formalize

- Safety (correctness)
- Liveness (progress)

for concurrent data structures and concurrent algorithms.

# Motivation

We need to formalize

- Safety (correctness)
- Liveness (progress)

for concurrent data structures and concurrent algorithms.

- How to define safety and liveness "universally"?
- How to reason about whole programs / software components in isolation?
- How to compose modules and correctness proofs?
- Are there "better" or "simpler" approaches?

# Concurrent objects

Goals:

- **Describe** object
- **Implement** object
- **Check** for errors

# Concurrent objects

Goals:

- **Describe** object
- **Implement** object
- **Check** for errors

**Check that implementation conforms to specification.**

# Concurrent objects

Goals:

- **Describe** object
- **Implement** object
- **Check** for errors

**Check that implementation conforms to specification.**

Let's focus on descriptions:

# Concurrent objects

Goals:

- **Describe** object
- **Implement** object
- **Check** for errors

**Check that implementation conforms to specification.**

Let's focus on descriptions:

- when implementation is **correct**
- the conditions under which it guarantees **progress**

# Sequential objects and specifications

Sequential object:

- State
  - Set of fields
- Methods
  - The only way to change state

# Sequential objects and specifications

Sequential object:

- State
  - Set of fields
- Methods
  - The only way to change state

Sequential specification:

- If (**precondition**)
  - the object is in such-and-such a state
  - before you call the method
- Then (**postcondition**)
  - the method will return a particular value
  - or throw a particular exception
- and (**postcondition**)
  - the object will be in some other state
  - when the method returns

# Sequential specifications

- Interactions among methods captured by side-effects on object state
  - State meaningful between method calls
- Documentation size linear in number of methods
  - Each method described in isolation
- Can add new methods
  - Without changing descriptions of old methods

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object might never be between method calls

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object might never be between method calls
  - in sequential specification we need to only define object state between method calls

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object might never be between method calls
  - in sequential specification we need to only define object state between method calls
- must characterize all possible interactions with concurrent calls

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object might never be between method calls
  - in sequential specification we need to only define object state between method calls
- must characterize all possible interactions with concurrent calls
  - in sequential specification we described methods in isolation

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object might never be between method calls
  - in sequential specification we need to only define object state between method calls
- must characterize all possible interactions with concurrent calls
  - in sequential specification we described methods in isolation
- everything can potentially interact with everything else

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object might never be between method calls
  - in sequential specification we need to only define object state between method calls
- must characterize all possible interactions with concurrent calls
  - in sequential specification we described methods in isolation
- everything can potentially interact with everything else
  - in sequential specification we could add new methods without affecting older methods

# Sequential vs concurrent

Key difference (concurrency):

- Method call is not an event
- Method call is an interval

Implications:

- method calls could overlap, object
  - in sequential specification we ne
- must characterize all possible interac
  - in sequential specification we de
- everything can potentially interact
  - in sequential specification we cc



# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - **Intuition behind consistency definitions**
  - Optional: formal definitions and composableity
- 4 Summary

# Consistency

What does it **mean** for a *concurrent* object to be correct?

- What is a concurrent FIFO queue?

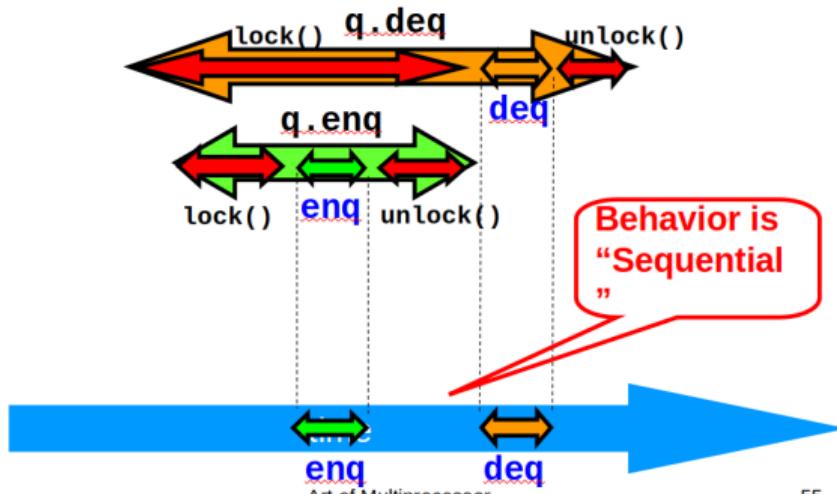
# Consistency

What does it **mean** for a *concurrent* object to be correct?

- What is a concurrent FIFO queue?
- FIFO means strict temporal order
- Concurrent means ambiguous temporal order

# Consistency

Lets capture the idea of describing the concurrent via the sequential



# Linearizability

Each method should

- “take effect”
- instantaneously
- between invocation and response events

# Linearizability

Each method should

- “take effect”
- instantaneously
- between invocation and response events

Object is correct if this “sequential” behavior is correct

# Linearizability

Each method should

- “take effect”
- instantaneously
- between invocation and response events

Object is correct if this “sequential” behavior is correct

Any such concurrent object is **Linearizable**

# Linearizability

Each method should

- “take effect”
- instantaneously
- between invocation and response events

Object is correct if this “sequential” behavior is correct

Any such concurrent object is **Linearizable**

Sounds like a property of an execution ...

# Linearizability

Each method should

- “take effect”
- instantaneously
- between invocation and response events

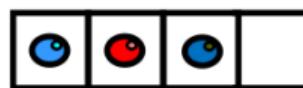
Object is correct if this “sequential” behavior is correct

Any such concurrent object is **Linearizable**

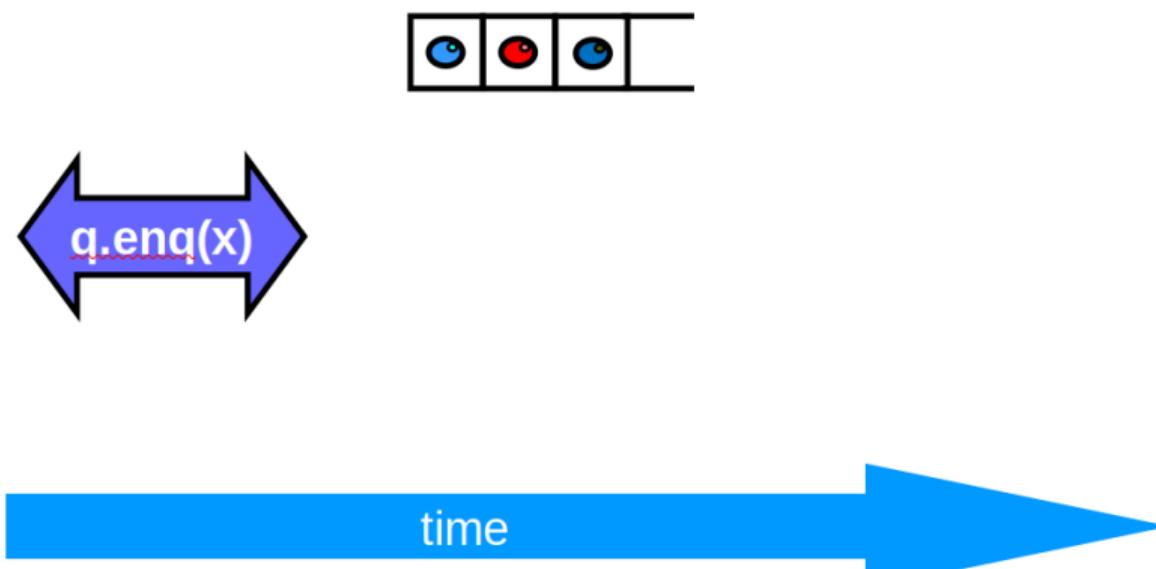
Sounds like a property of an execution ...

A **linearizable object**: one all of whose possible executions are linearizable

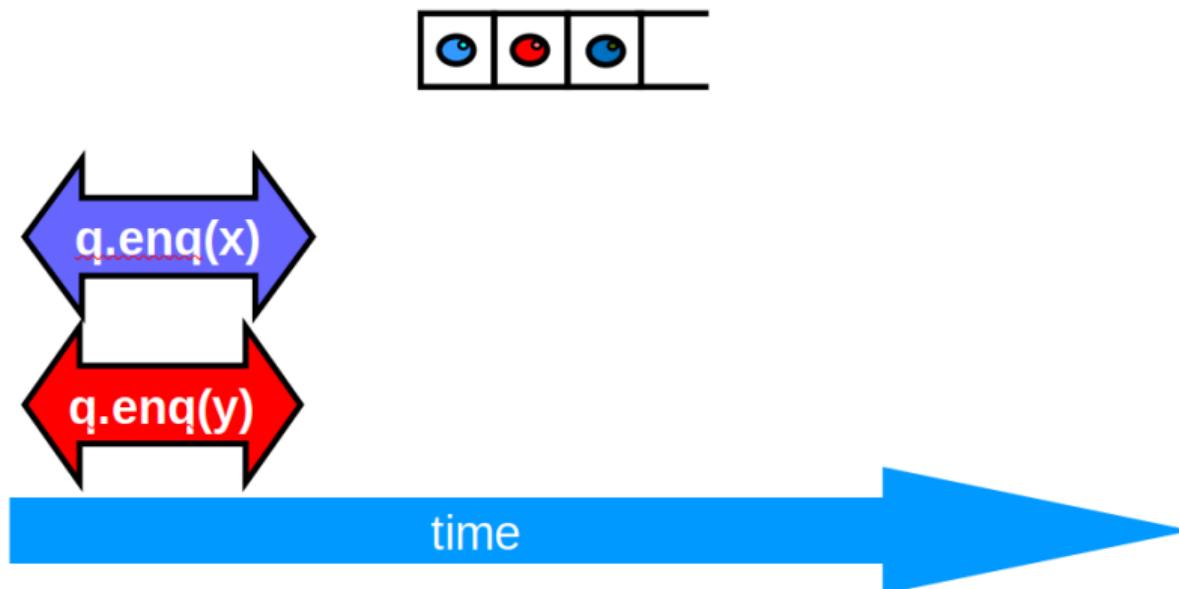
# Examples



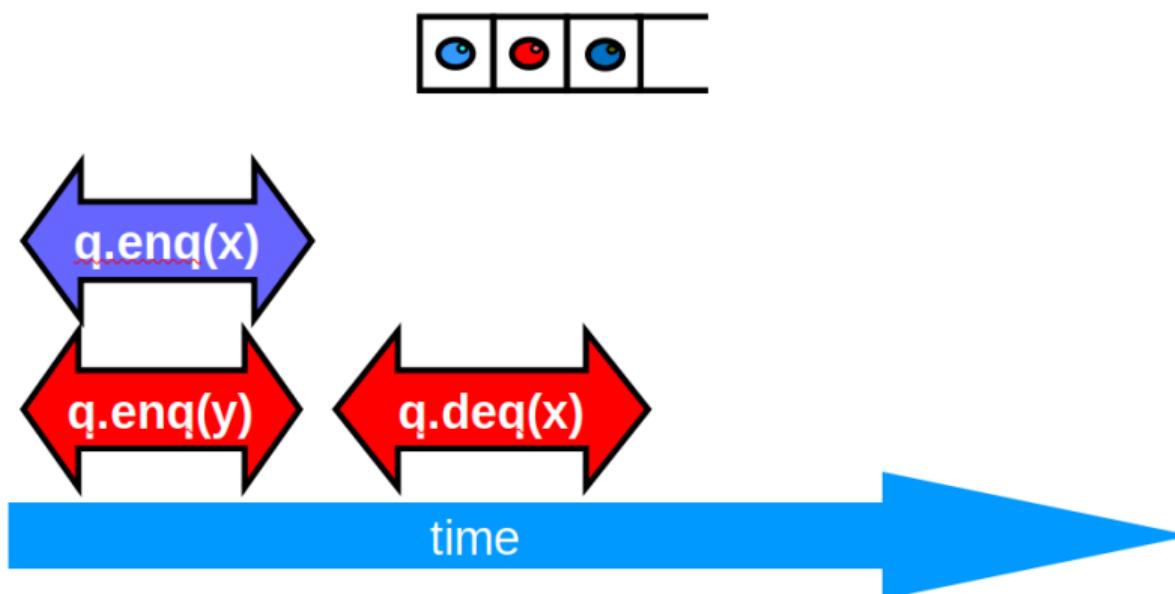
# Examples



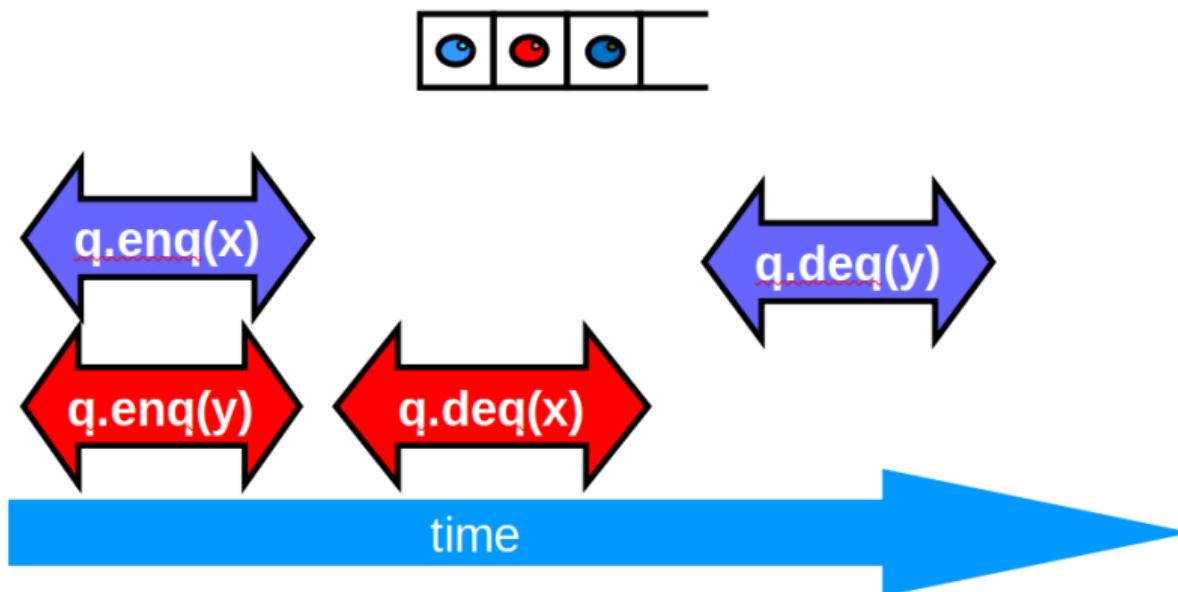
# Examples



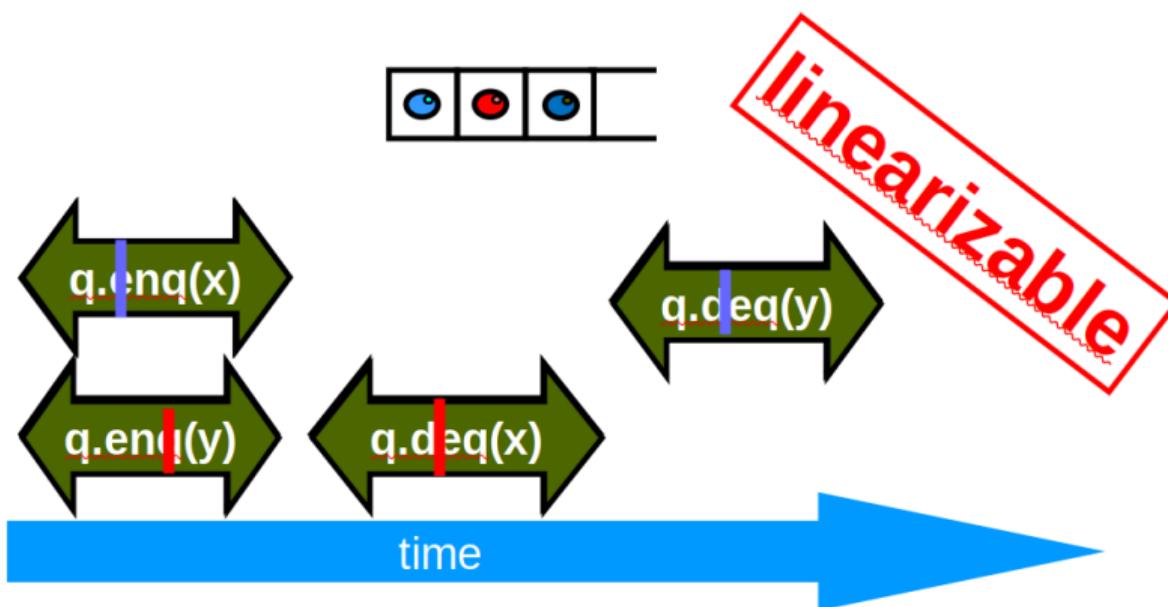
## Examples



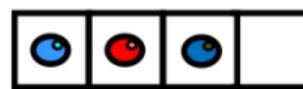
## Examples



## Examples

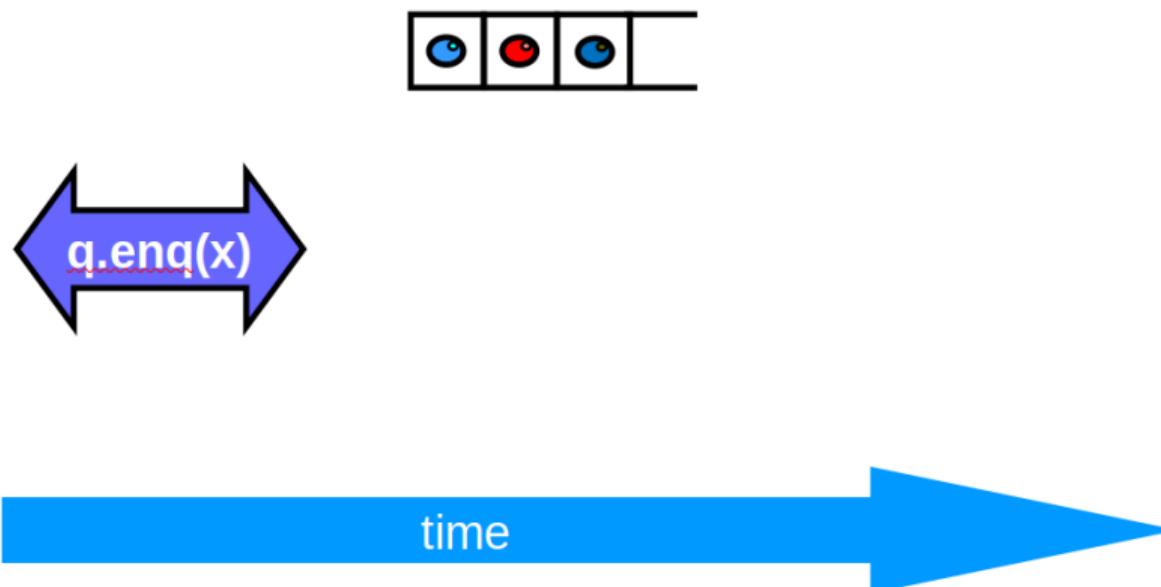


# Examples

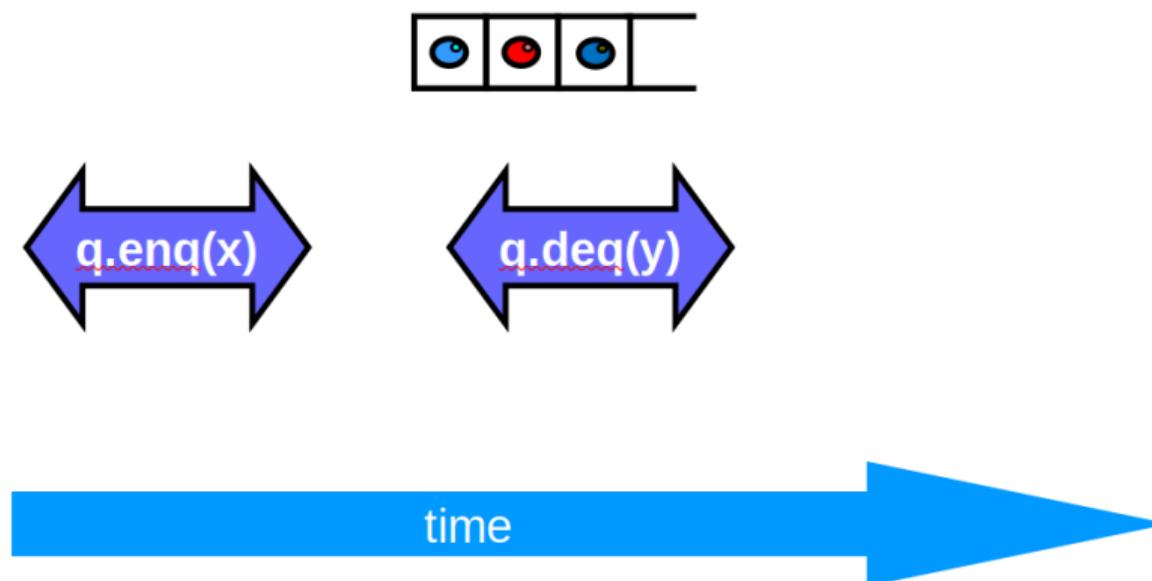


time

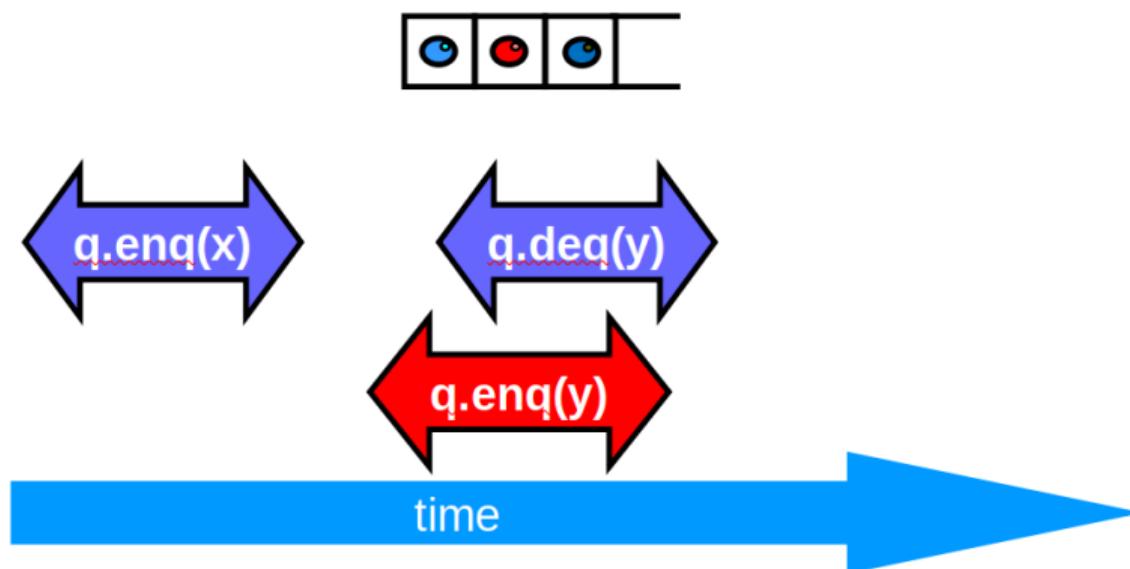
# Examples



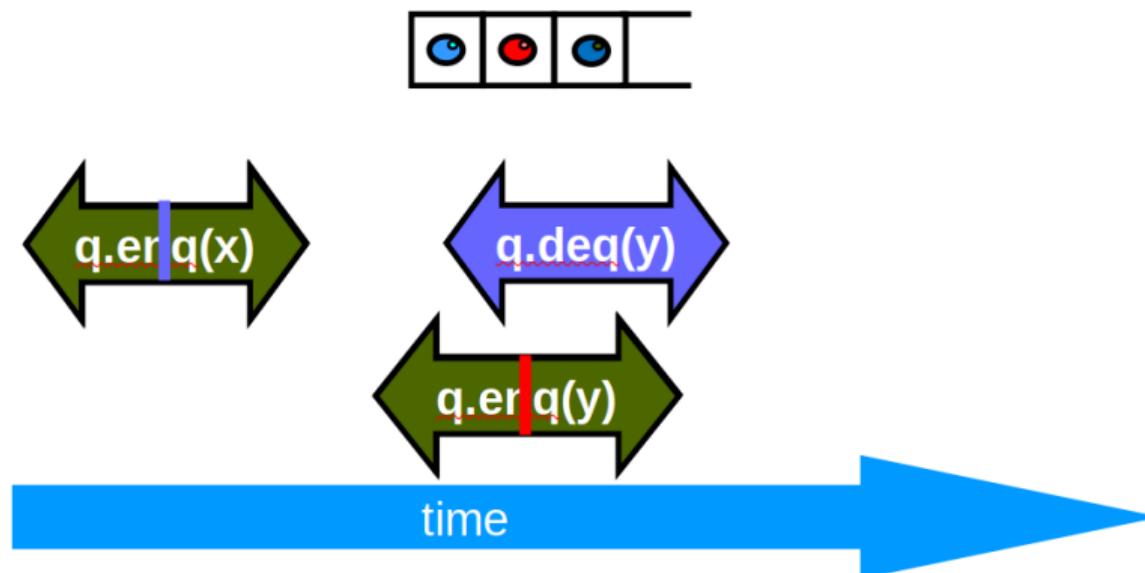
## Examples



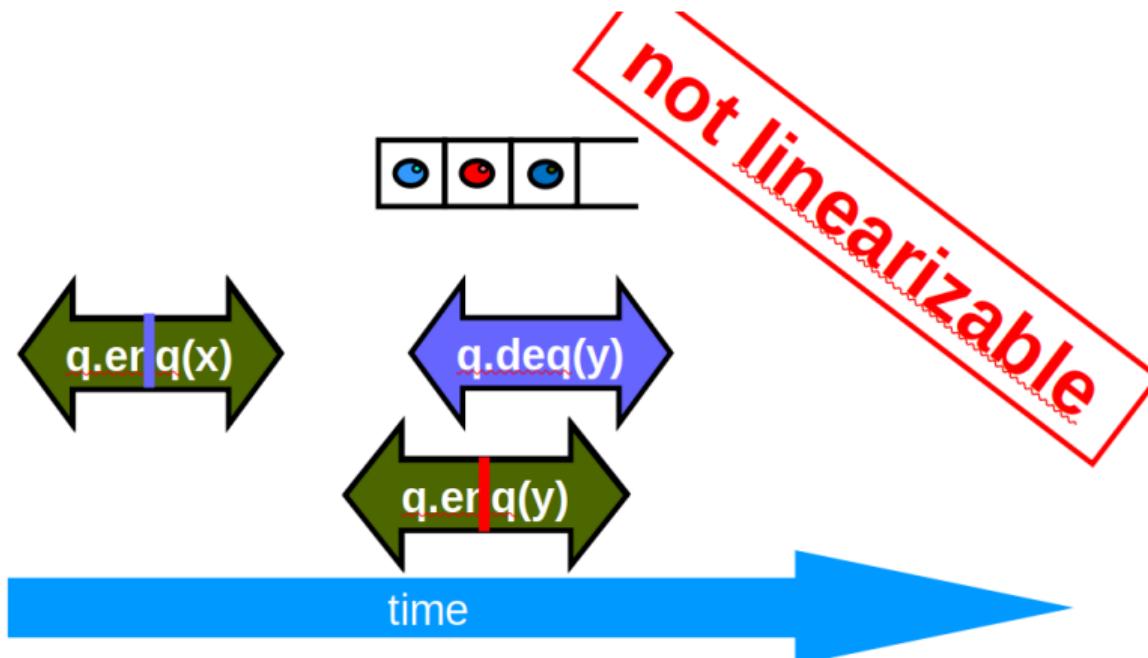
## Examples



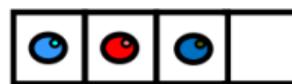
## Examples



## Examples

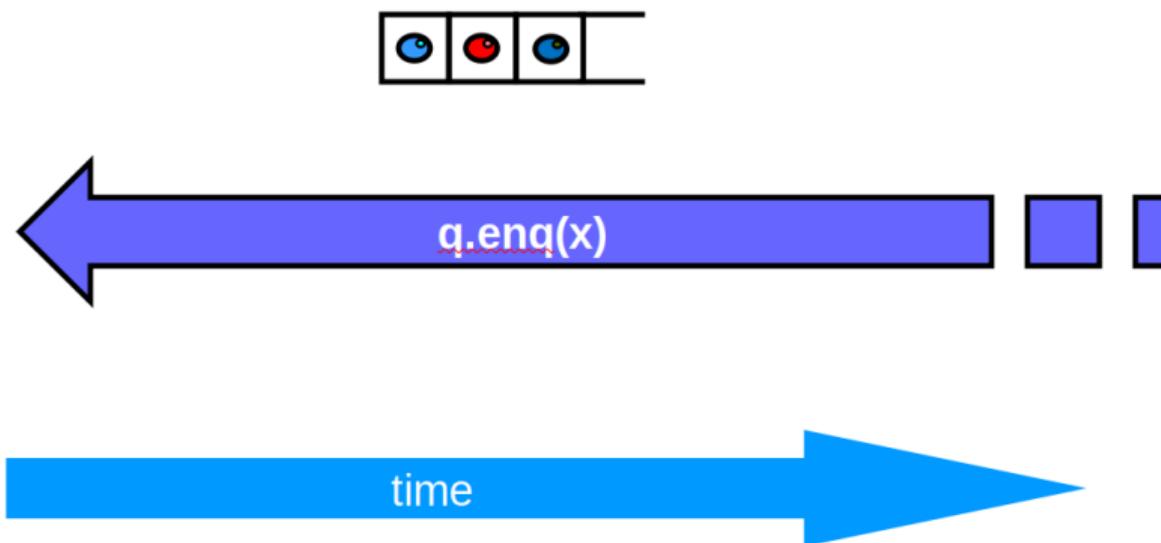


# Examples

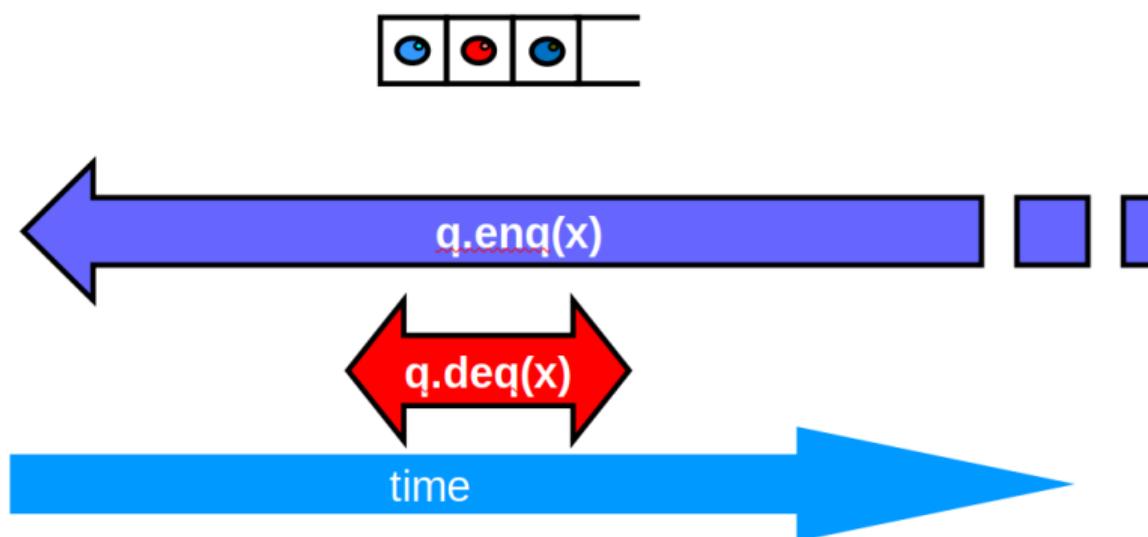


time

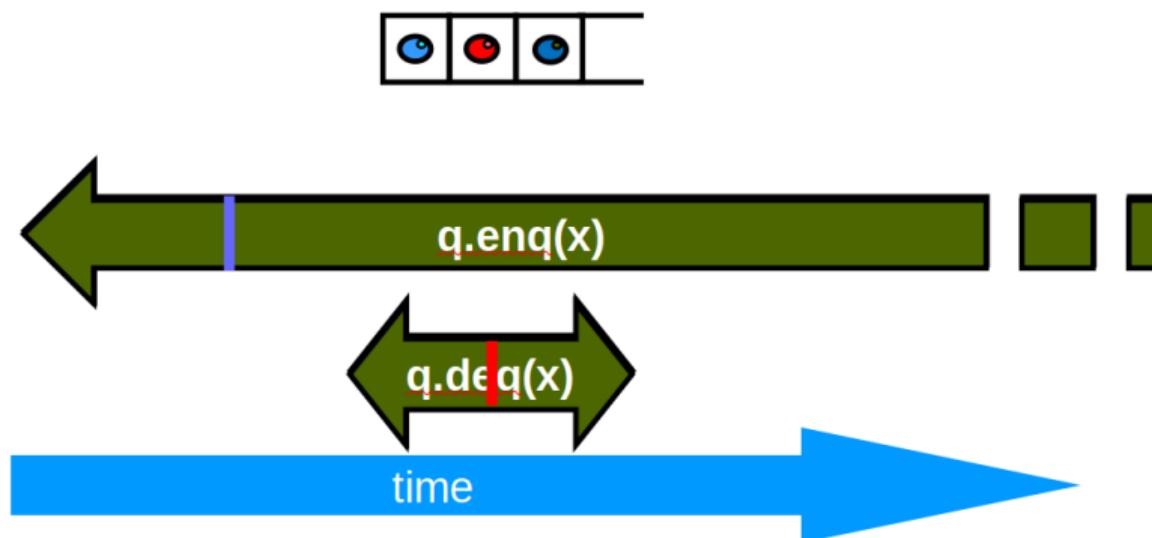
# Examples



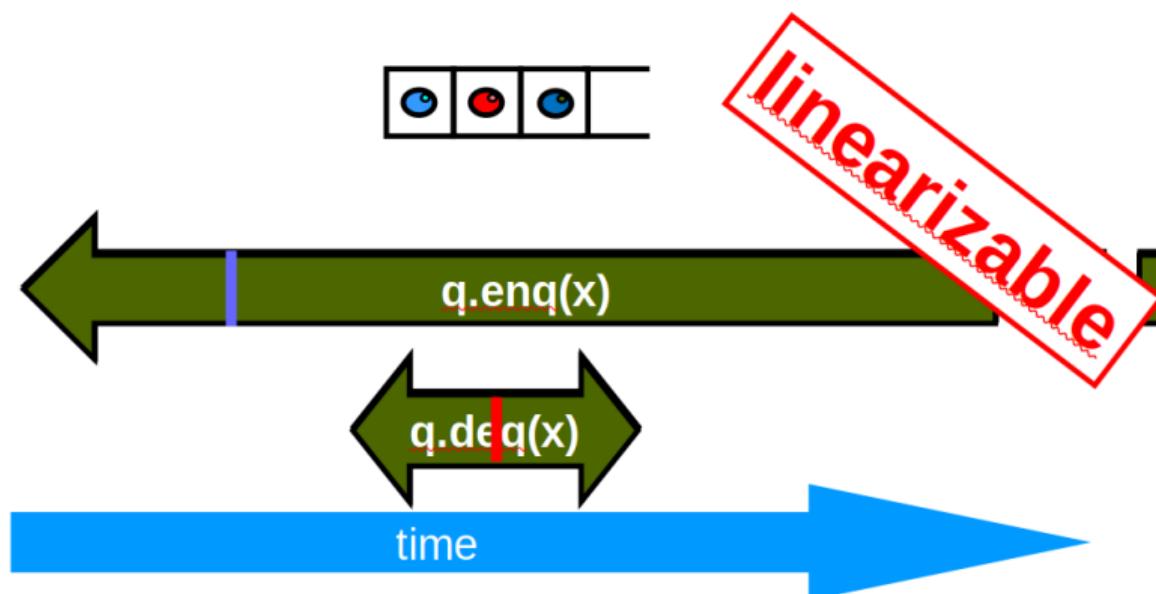
## Examples



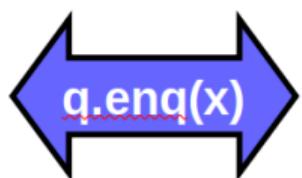
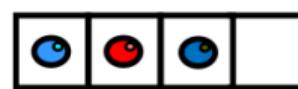
## Examples



## Examples

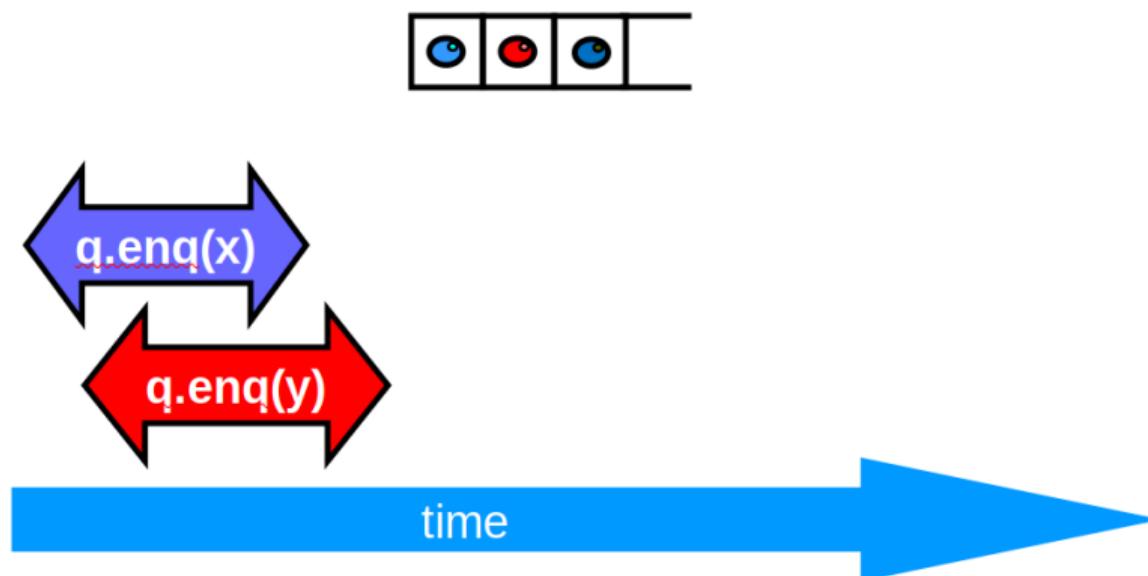


## Examples

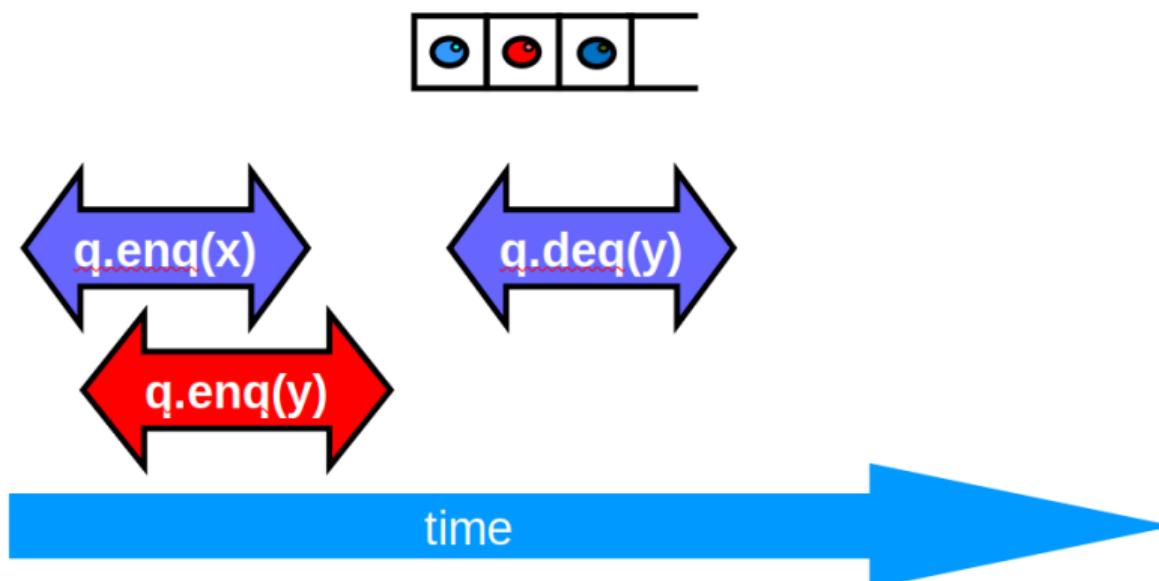


time

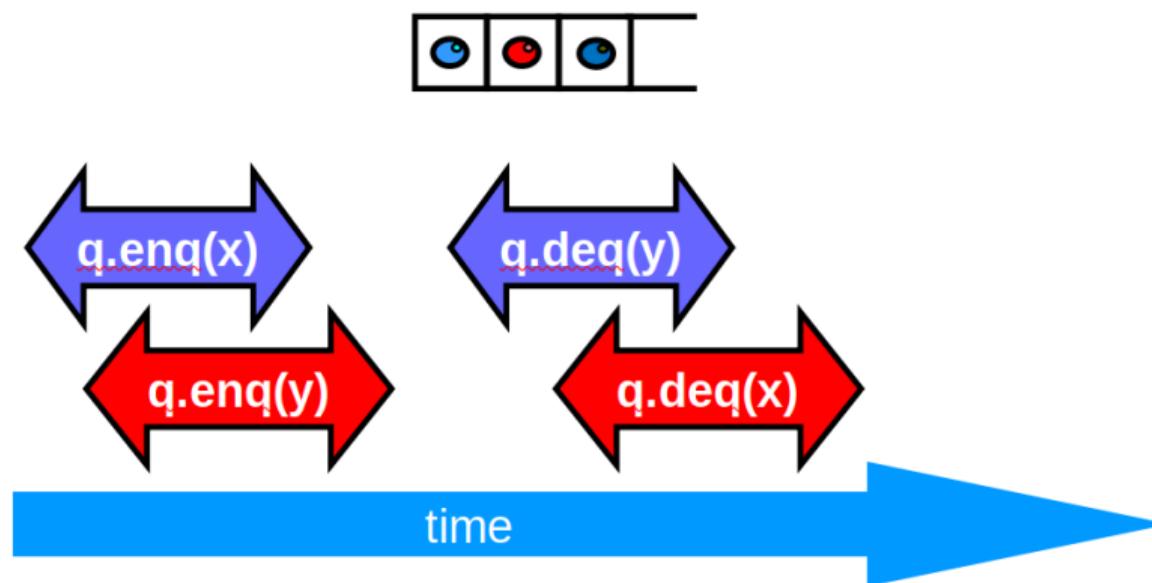
## Examples



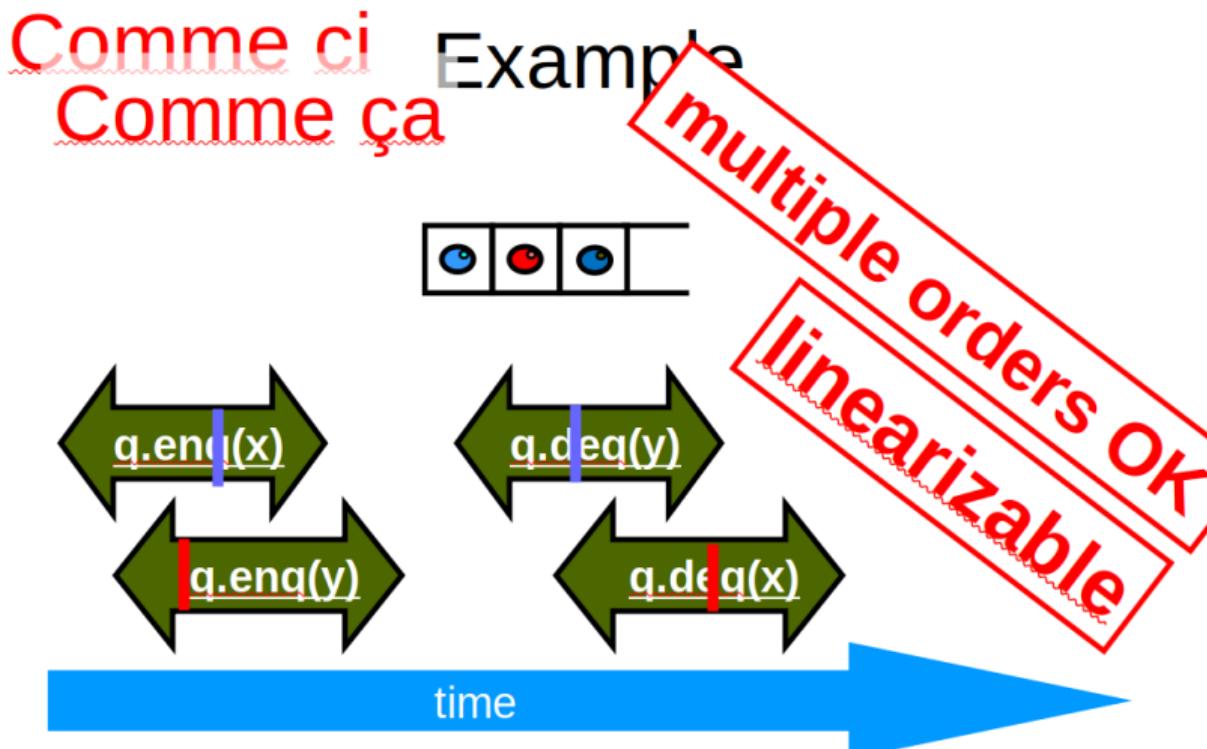
## Examples



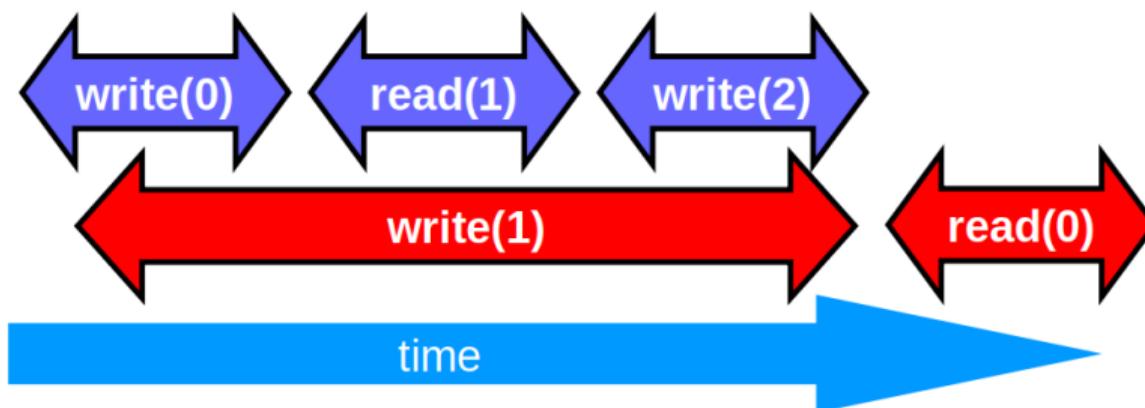
## Examples



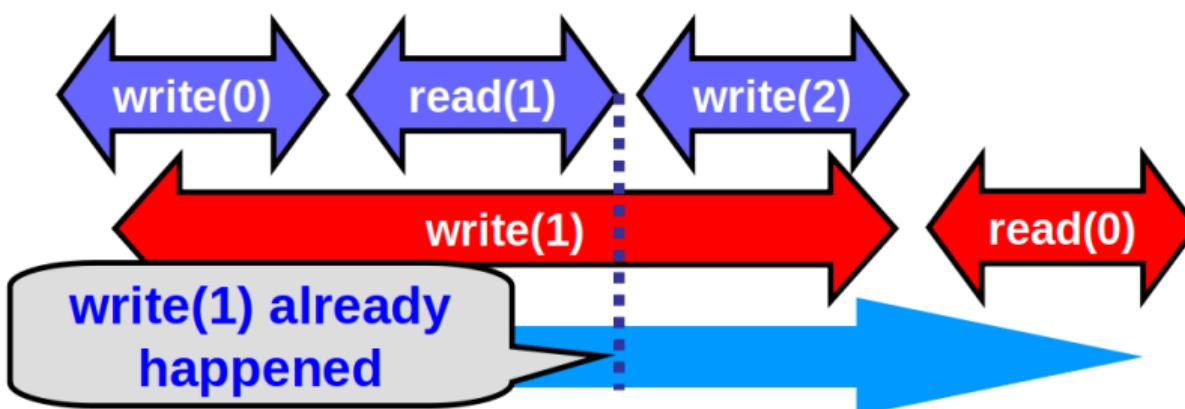
## Examples



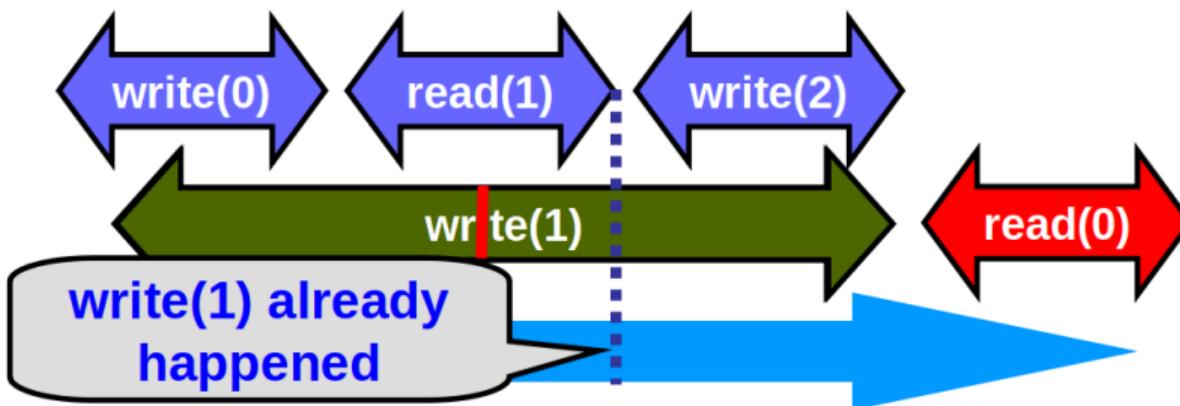
## Examples



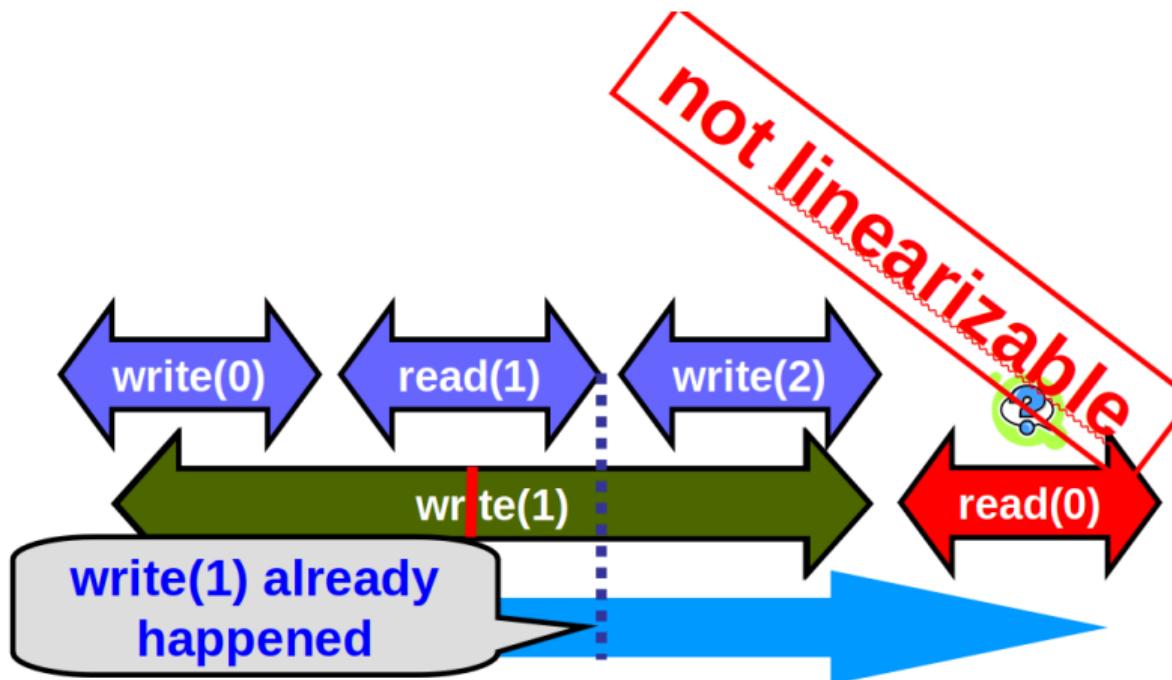
## Examples



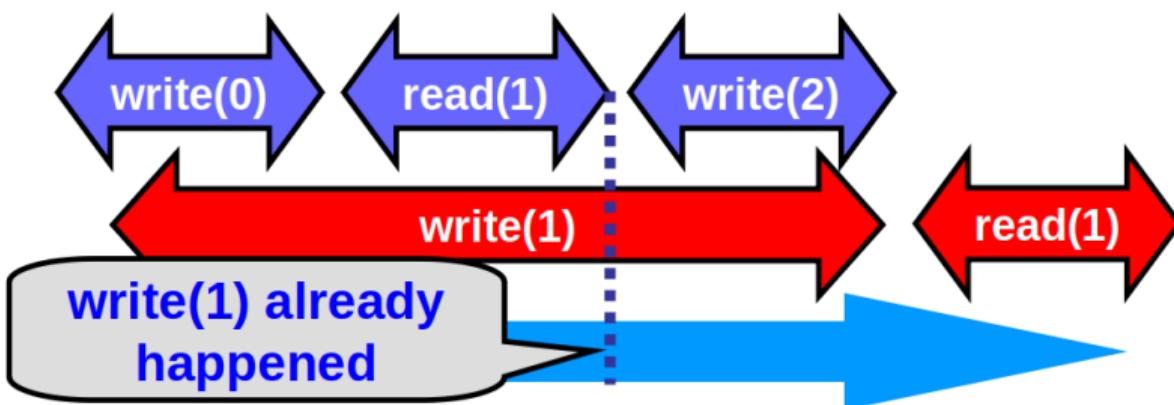
## Examples



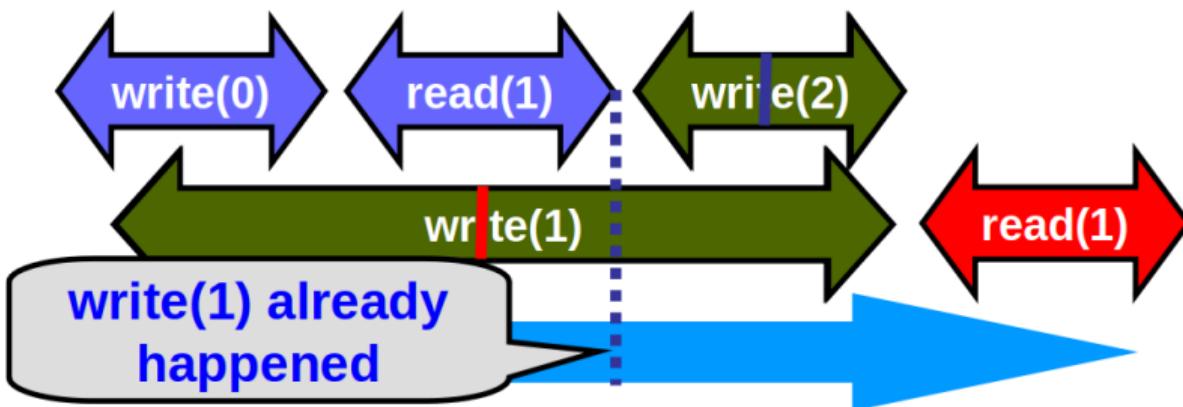
## Examples



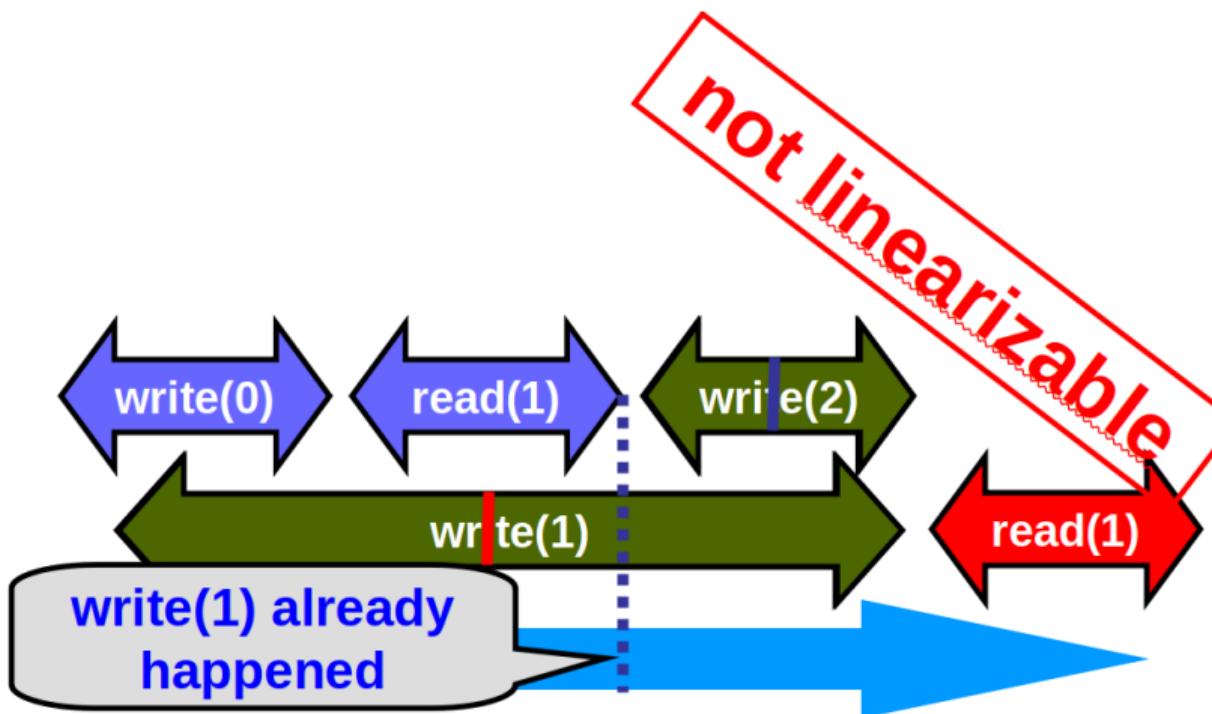
## Examples



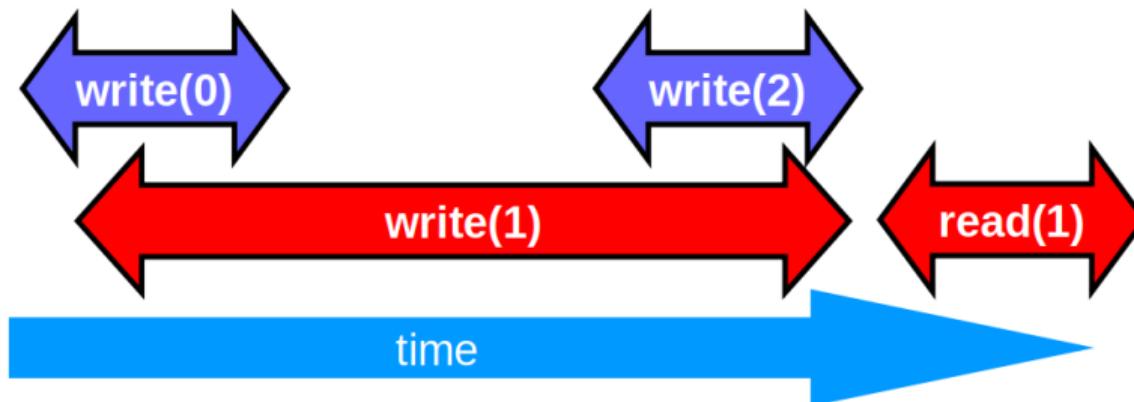
## Examples



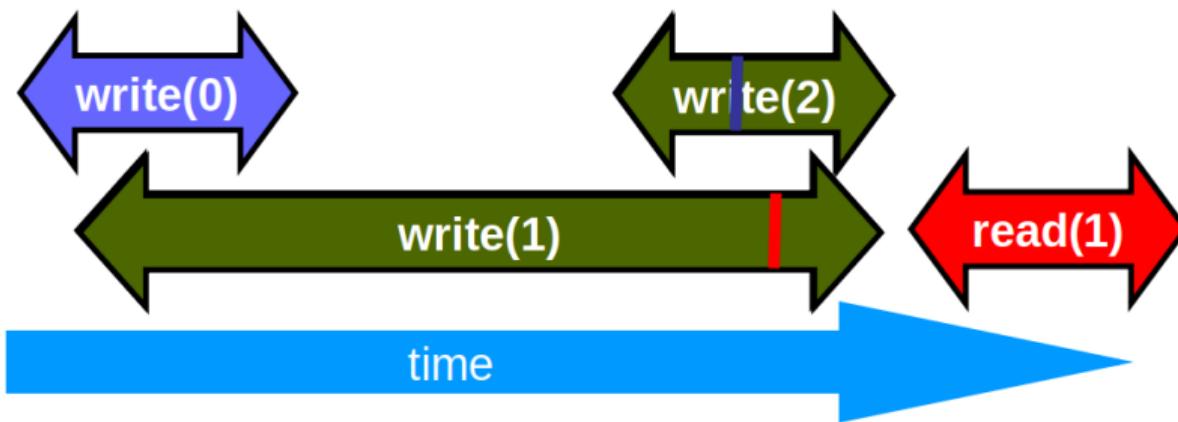
## Examples



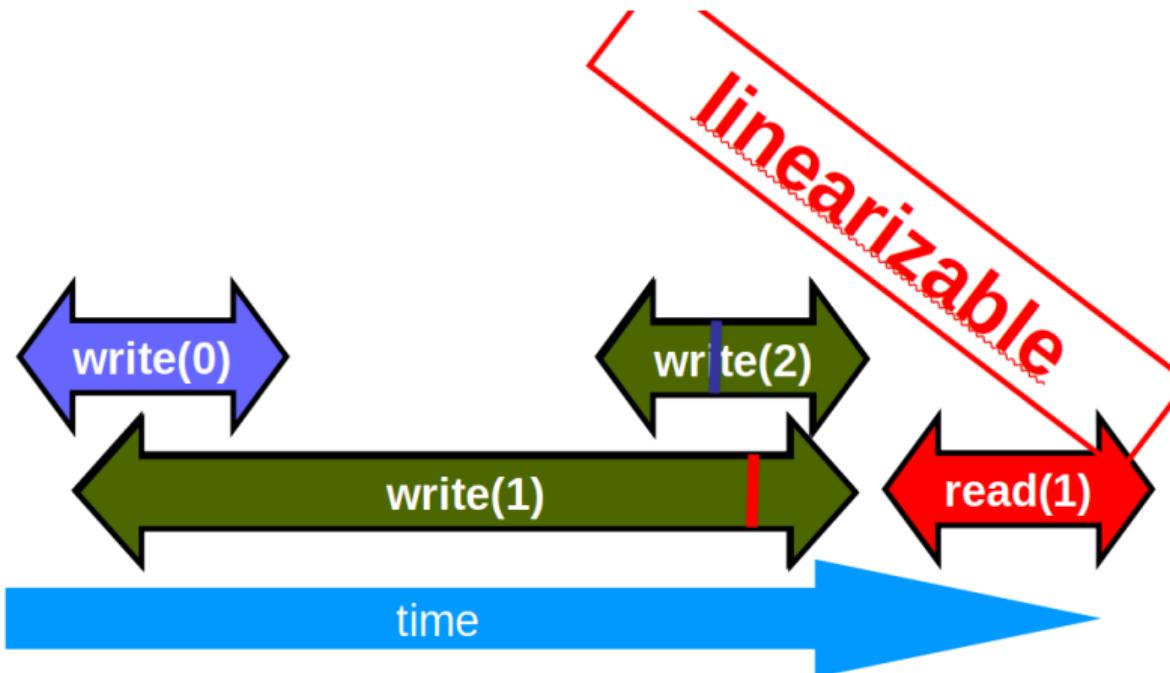
## Examples



## Examples



## Examples



# Linearizability: summary

Each method should

- “take effect”
- instantaneously
- between invocation and response events

Linearization point: when method “commit changes to object state”

- Could be different code points for different executions!

Linearizability

- requires to analyze executions, not just methods
- compositional – result of composing linearizable objects is linearizable<sup>2</sup>

---

<sup>2</sup>Optional homework: Theorem 3.6.1 in "The Art of Multiprocessor Programming"

# Linearizability is not the only one

There exists:

# Linearizability is not the only one

There exists:

- Sequential consistency
  - Provides other guarantees, harder to reason (no linearization points)
  - **Not** compositional
  - Hardware designers love to use weaker variations of it (Lecture 9 will be hard)

# Linearizability is not the only one

There exists:

- Sequential consistency
  - Provides other guarantees, harder to reason (no linearization points)
  - **Not** compositional
  - Hardware designers love to use weaker variations of it (Lecture 9 will be hard)
- Quiescent consistency
  - One of the weakest (yet understandable) consistencies
  - Compositional
  - Complicated performance-oriented concurrent data structures may use it to formalize possible and impossible results

# Linearizability is not the only one

There exists:

- Sequential consistency
  - Provides other guarantees, harder to reason (no linearization points)
  - **Not** compositional
  - Hardware designers love to use weaker variations of it (Lecture 9 will be hard)
- Quiescent consistency
  - One of the weakest (yet understandable) consistencies
  - Compositional
  - Complicated performance-oriented concurrent data structures may use it to formalize possible and impossible results

Linearizability is useful yet could be not applicable (by design<sup>3</sup> or by implementation<sup>4</sup>)

---

<sup>3</sup>RCU: advantages and disadvantages <https://en.wikipedia.org/wiki/Read-copy-update>

<sup>4</sup>ConcurrentLinkedDeque is non-linearizable <https://bugs.openjdk.org/browse/JDK-8256833>

## Intuition behind consistency definitions: takeaways

- **Describe** object first, then implement and test it
- Sequential specifications totally rock
  - Easy to understand and document
  - Modular and extensible
- Concurrent methods are not instantaneous and could overlap

## Intuition behind consistency definitions: takeaways

- **Describe** object first, then implement and test it
- Sequential specifications totally rock
  - Easy to understand and document
  - Modular and extensible
- Concurrent methods are not instantaneous and could overlap

Solution 0: use locks and sequential specifications

## Intuition behind consistency definitions: takeaways

- **Describe** object first, then implement and test it
- Sequential specifications totally rock
  - Easy to understand and document
  - Modular and extensible
- Concurrent methods are not instantaneous and could overlap

Solution 0: use locks and sequential specifications

Solution 1: mimic to sequential specification

- Points in code where method "committed changes" : linearization points
- Every concurrent execution have order among object-specific operations: linearizable object
- Provide specifications based on linearizability

## Intuition behind consistency definitions: takeaways

- **Describe** object first, then implement and test it
- Sequential specifications totally rock
  - Easy to understand and document
  - Modular and extensible
- Concurrent methods are not instantaneous and could overlap

Solution 0: use locks and sequential specifications

Solution 1: mimic to sequential specification

- Points in code where method "committed changes" : linearization points
- Every concurrent execution have order among object-specific operations: linearizable object
- Provide specifications based on linearizability

Solution 2: use more fragile and complicated consistency definitions

## Intuition behind consistency definitions: takeaways

- **Describe** object first, then implement and test it
- Sequential specifications totally rock
  - Easy to understand and document
  - Modular and extensible
- Concurrent methods are not instantaneous and could overlap

Solution 0: use locks and sequential specifications

Solution 1: mimic to sequential specification

- Points in code where method "committed changes" : linearization points
- Every concurrent execution have order among object-specific operations: linearizable object
- Provide specifications based on linearizability

Solution 2: use more fragile and complicated consistency definitions

Remember to check if your approach is compositional.

# Lecture plan

- 1 Time, events, intervals, precedence
- 2 Mutual exclusion
  - 2 threads
  - N threads
  - Lower bounds on the Number of Locations
  - Optional: fairness
  - Mutual exclusion beyond read-write memory cells
- 3 Concurrent objects
  - Sequential and concurrent objects
  - Intuition behind consistency definitions
  - **Optional: formal definitions and composability**
- 4 Summary

## Formal definitions

Optional homework: slides 95-176 from chapter\_03.ppt

Optional homework: Chapter 3 "Concurrent Objects" up to section 3.6 "Formal Definitions"  
(pages 45 - 58)

Key insights:

- quiescent consistency
- sequential consistency
- linearizability
- formal definitions: history, projections
- compositional linearizability
- nonblocking property

# Summary

First steps towards formalization of concurrent execution

- Timeline, Event, Interval, Precedence

Rigorously prove mutual exclusion properties:

- Mutual exclusion, Deadlock-freedom, Starvation-freedom

Non-trivial results for mutual exclusion

- Peterson's algorithm
- FilterLock
- Lower bounds on the number of locations

Formalizing allowed behaviour of concurrent object:

- Sequential object, sequential specification
- Concurrent object, consistency
- Linearizability, linearization points
- Non-linearizable executions

## Summary: homework

Your second critical task will be related to the theoretical concurrency.

Maybe you would like to re-read slides and check out companion slides for "The Art of Multiprocessor Programming"